

CSE 31


Computer Organization

Lecture 13 – Instruction Format (2)

Announcement

- ▶ Project #1
 - Due at 11:59pm on 3/22, Friday (no more late submission)
 - You must demo your submission to your TA during week of 4/1, in lab.
- ▶ Lab #7 this week
 - Due after Spring Break
 - Demo your lab within 7 days after due dates
- ▶ HW #4 in CatCourses
 - Due Monday (4/1) at 11:59pm
- ▶ Reading assignment
 - Chapter 1.6, 6.1-6.3 of zyBooks
 - Make sure to do the Participation Activities
 - Due Wednesday (4/3) at 11:59pm

Levels of Representation (abstractions)



Instruction Formats

- ▶ **3 formats**
- ▶ **I-format**: used for instructions with immediates, **lw** and **sw** (since offset counts as an immediate), and branches (**beq** and **bne**),
 - (but not the shift instructions; later)
- ▶ **J-format**: used for **j** and **jal**
- ▶ **R-format**: used for all other instructions
- ▶ **Why 3 different formats?**
 - It will soon become clear why the instructions have been partitioned in this way.

R-Format Instructions (1/5)

- ▶ Define “**fields**” of the following number of bits each: $6 + 5 + 5 + 5 + 5 + 6 = 32$

6	5	5	5	5	6
---	---	---	---	---	---

- ▶ For simplicity, each field has a name:

opcode	rs	rt	rd	shamt	funct
--------	----	----	----	-------	-------

- ▶ **Important:** On these slides and in book, each field is viewed as a 5- or 6-bit unsigned integer, not as part of a 32-bit integer.
 - Consequence: 5-bit fields can represent any number 0-31, while 6-bit fields can represent any number 0-63.

R-Format Instructions (2/5)

- ▶ What do these field integer values tell us?
 - **opcode**: partially specifies what instruction it is
 - **Note: This number is equal to 0 for all R-Format instructions.**
 - **funct**: combined with `opcode`, this number exactly specifies the instruction
- ▶ Question: Why aren't `opcode` and `funct` a single 12-bit field?
 - We'll answer this later.

R-Format Instructions (3/5)

- ▶ More fields:
 - rs (Source Register): **generally** used to specify register containing **first operand**
 - rt (Target Register): **generally** used to specify register containing **second operand** (note that name is misleading)
 - rd (Destination Register): **generally** used to specify register which will **receive result** of computation

R-Format Instructions (4/5)

- ▶ Notes about register fields:
 - Do we need more than 5 bits?
 - Each register field is exactly 5 bits, which means that it can specify any unsigned integer in the range 0-31. Each of these fields specifies one of the 32 registers by number.
 - The word “generally” was used because there are exceptions
 - **mult** and **div** have nothing important in the **rd** field since the dest registers are **hi** and **lo**
 - **mfhi** and **mflo** have nothing important in the **rs** and **rt** fields since the source is determined by the instruction.

R-Format Instructions (5/5)

- ▶ Final field:
 - shamt: This field contains the amount a shift instruction will shift by. Shifting a 32-bit word by more than 31 is useless, so this field is only 5 bits (so it can represent the numbers 0-31).
 - **This field is set to 0 in all but the shift instructions.**
- ▶ For a detailed description of field usage for each instruction, see MIPS_Reference_Sheet.pdf in CatCourses (You will be provided in all exams)

R-Format Example (1/2)

► MIPS Instruction:

add \$t0,\$t1,\$t2

add \$8,\$9,\$10

opcode = 0 (look up in table)

funct = 32 (look up in table)

rd = 8 (destination)

rs = 9 (first operand)

rt = 10 (second operand)

shamt = 0 (not a shift)

MIPS Reference Data

CORE INSTRUCTION SET

NAME, MNEMONIC	FOR- MAT	OPERATION (in Verilog)
Add	add R	$R[rd] = R[rs] + R[rt]$

①



OPCODE

/FUNCT

(Hex)

(1) 0 20_{hex}

R-Format Example (2/2)

► MIPS Instruction:

add \$8, \$9, \$10

Decimal number per field representation:

0	9	10	8	0	32
---	---	----	---	---	----

Binary number per field representation:

000000	01001	01010	01000	00000	100000
--------	-------	-------	-------	-------	--------

hex representation:

012A 4020_{hex}

decimal representation:

19,546,144_{ten}

Called a Machine Language Instruction

hex

I-Format Instructions (1/4)

- ▶ What about instructions with immediates?
 - 5-bit field only represents numbers up to the value 31: immediates may be much larger than this
 - Ideally, MIPS would have only one instruction format (for simplicity): unfortunately, we need to compromise
- ▶ Define new instruction format that is partially consistent with R-format:
 - First notice that, if instruction has immediate, then it uses at most 2 registers.

I-Format Instructions (2/4)

- ▶ Define “fields” of the following number of bits each: $6 + 5 + 5 + 16 = 32$ bits

6	5	5	16
---	---	---	----

- Again, each field has a name:

opcode	rs	rt	immediate
--------	----	----	-----------

- **Key Concept:** Only one field is inconsistent with R-format. Most importantly, `opcode` is still in same location.

I-Format Instructions (3/4)

- ▶ What do these fields mean?
 - **opcode**: same as before except that, since there's no `funct` field, `opcode` uniquely specifies an instruction in I-format
 - This also answers question of why R-format has two 6-bit fields to identify instruction instead of a single 12-bit field: in order to be consistent as possible with other formats while leaving as much space as possible for immediate field.
 - **rs**: specifies a register operand (if there is one)
 - **rt**: specifies register which will receive result of computation (this is why it's called the **target** register "`rt`") or other operand for some instructions.

I-Format Instructions (4/4)

▶ The Immediate Field:

- **addi**, **slti**, **sltiu**, the immediate is **sign-extended** to 32 bits. Thus, it's treated as a signed integer.
- 16 bits → can be used to represent immediate up to 2^{16} different values
- This is large enough to handle the offset in a typical **lw** or **sw**, plus a vast majority of values that will be used in the **slti** instruction.
- We'll see what to do when the number is too big later

I-Format Example (1/2)

- ▶ MIPS Instruction:

addi **\$s5, \$s6, -50**

addi **\$21, \$22, -50**

opcode = 8 (look up in table in book)

rs = 22 (register containing operand)

rt = 21 (target register)


immediate = -50 (by default, this is decimal)

Add Immediate

addi

I $R[rt] = R[rs] + \text{SignExtImm}$

(1,2)



I-Format Example (2/2)

- ▶ MIPS Instruction:

`addi $21, $22, -50`

Decimal/field representation:

8	22	21	-50
---	----	----	-----

Binary/field representation:

001000	10110	10101	1111111111001110
--------	-------	-------	------------------

hexadecimal representation: `22D5 FFCEhex`

decimal representation: `584, 449, 998ten`

Quiz

Which instruction has same representation as 35_{ten} ?

- a) add \$0, \$0, \$0
- b) subu \$s0,\$s0,\$s0
- c) lw \$0, 0(\$0)
- d) addi \$0, \$0, 35
- e) subu \$0, \$0, \$0

opcode	rs	rt	rd	shamt	funct
opcode	rs	rt	rd	shamt	funct
opcode	rs	rt	offset		
opcode	rs	rt	immediate		
opcode	rs	rt	rd	shamt	funct

Registers numbers and names:

0: \$0, .. 8: \$t0, 9: \$t1, ..15: \$t7, 16: \$s0, 17: \$s1, .. 23: \$s7

Opcodes and function fields (if necessary)

add: opcode = 0, funct = 32

subu: opcode = 0, funct = 35

addi: opcode = 8

lw: opcode = 35

Quiz

Which instruction has same representation as 35_{ten} ?

- a) add \$0, \$0, \$0
- b) subu \$s0,\$s0,\$s0
- c) lw \$0, 0(\$0)
- d) addi \$0, \$0, 35
- e) subu \$0, \$0, \$0**

opcode	rs	rt	rd	shamt	funct
opcode	rs	rt	rd	shamt	funct
opcode	rs	rt	offset		
opcode	rs	rt	immediate		
opcode	rs	rt	rd	shamt	funct

Registers numbers and names:

0: \$0, .. 8: \$t0, 9:\$t1, ..15:

\$t7, 16: \$s0, 17: \$s1, .. 23: \$s7

Opcodes and function fields (if necessary)

add: opcode = 0, funct = 32

subu: opcode = 0, funct = 35

addi: opcode = 8

lw: opcode = 35

I-Format Problems (0/3)

▶ Problem 0: Unsigned # sign-extended?

- `addiu`, `sltiu`, **sign-extends** immediates to 32 bits. Thus, # is a “signed” integer.

▶ Rationale

- `addiu` so that can add w/out overflow
 - See K&R pp. 230, 305
- `sltiu` suffers so that we can have easy HW
 - Does this mean we’ll get wrong answers?
 - Nope, it means assembler has to handle any unsigned immediate $2^{15} \leq n < 2^{16}$ (i.e., with a 1 in the 15th bit and 0s in the upper 2 bytes) as it does for numbers that are too large.

I-Format Problem (1/3)

► Problem:

- Chances are that `addi`, `lw`, `sw` and `slti` will use immediates small enough to fit in the immediate field.
- ...but what if it's too big?
 - `addi $s0, $s1, 40000`
 - We need a way to deal with a 32-bit immediate in any I-format instruction.

I-Format Problem (2/3)

▶ Solution to Problem:

- Handle it in software + new instruction
- Don't change the current instructions: instead, add a new instruction to help out

▶ New instruction:

`lui register, immediate`

- stands for ***Load Upper Immediate***
- takes 16-bit immediate and puts these bits in the upper half (high order half) of the register
- sets lower half to 0s

I-Format Problems (3/3)

► Solution to Problem (continued):

- So how does `lui` help us?
- Example:

```
addiu $t0, $t0, 0xABABCD
```

...becomes

```
lui $at 0xABAB
```

```
ori $at, $at, 0xCD
```

```
addu $t0, $t0, $at
```

R-format!

- Now each I-format instruction has only a 16-bit immediate.
- Wouldn't it be nice if the assembler would do this for us automatically? (later)

Branches: PC-Relative Addressing (1/5)

- ▶ Use I-Format

opcode	rs	rt	immediate
--------	----	----	-----------

- ▶ `opcode` specifies `beq` versus `bne`
- ▶ `rs` and `rt` specify registers to compare
- ▶ What can `immediate` specify?
 - `immediate` is only 16 bits
 - PC (Program Counter) has byte address of current instruction being executed;
 - 32-bit pointer to memory
 - So `immediate` cannot specify entire address to branch to.

Branches: PC-Relative Addressing (2/5)

- ▶ How do we typically use branches?
 - Answer: `if-else`, `while`, `for`
 - Loops are generally small: usually up to 50 instructions
 - Function calls and unconditional jumps are done using jump instructions (`j` and `jal`), not the branches.
- ▶ Conclusion: may want to branch to anywhere in memory, but a branch often changes **PC** by a small amount

Branches: PC-Relative Addressing (3/5)

- ▶ Solution to branches in a 32-bit instruction:
 - PC-Relative Addressing
- ▶ Let the 16-bit immediate field be a signed two's complement integer to be added to the PC if we take the branch.
- ▶ Now we can branch $\pm 2^{15}$ bytes from the PC, which should be enough to cover almost any loop.
- ▶ Any ideas to further optimize this?

Branches: PC-Relative Addressing (4/5)

- ▶ Note: Instructions are words, so they're word aligned (byte address is always a multiple of 4, which means it ends with 00 in binary).
 - So the number of bytes to add to the PC will always be a multiple of 4.
 - So specify the `immediate` in words.
- ▶ Now, we can branch $\pm 2^{15}$ words from the PC (or $\pm 2^{17}$ bytes), so we can handle loops 4 times as large.

Branches: PC-Relative Addressing (5/5)

► Branch Calculation:

- If we **don't** take the branch:

$PC = PC + 4 = \text{byte address of next instruction}$

- If we **do** take the branch:

$PC = (PC + 4) + (\text{immediate} * 4)$

- Observations

- Immediate field specifies the number of words to jump, which is simply the number of instructions to jump.
- Immediate field can be positive or negative.
- Due to hardware, add immediate to (PC+4), not to PC; will be clearer (why? later in course)

Branch Example (1/3)

▶ MIPS Code:

```
Loop:  beq    $9, $0, End
        addu   $8, $8, $10
        addiu  $9, $9, -1
        j      Loop
```

End:

▶ beq branch is I-Format:

opcode = 4 (look up in table)

rs = 9 (first operand)

rt = 0 (second operand)

immediate = ??? (Do we need to know the address of End?)

Branch Example (2/3)

▶ MIPS Code:

```
Loop: beq    $9, $0, End ← PC
      addu   $8, $8, $10 ← PC + 4
      addiu  $9, $9, -1
      j      Loop
End:
```

▶ immediate Field:

- Number of **instructions** to add to (or subtract from) the PC, starting at the instruction **following** the branch.
- In **beq** case, **immediate** = 3

Branch Example (3/3)

► MIPS Code:

```
Loop:  beq    $9, $0, End
        addu   $8, $8, $10
        addiu  $9, $9, -1
        j      Loop
```

End:

decimal representation:

4	9	0	3
---	---	---	---

binary representation:

000100	01001	00000	000000000000000011
--------	-------	-------	--------------------

Questions on PC-addressing

- ▶ Does the value in branch field change if we move the code?
- ▶ What do we do if destination is $> 2^{15}$ instructions away from branch?
- ▶ Why do we need different addressing modes (different ways of describing a memory address)? Why not just one?

Summary

- ▶ Simplifying MIPS: Define instructions to be same size as data word (one word) so that they can use the same memory (compiler can use lw and sw).
- ▶ Computer actually stores programs as a series of these 32-bit numbers.
- ▶ **MIPS Machine Language Instruction:**
32 bits representing a single instruction

R	opcode	rs	rt	rd	shamt	funct
I	opcode	rs	rt	immediate		