

CSE 31

Computer Organization

Lecture 20 – Cache (2)

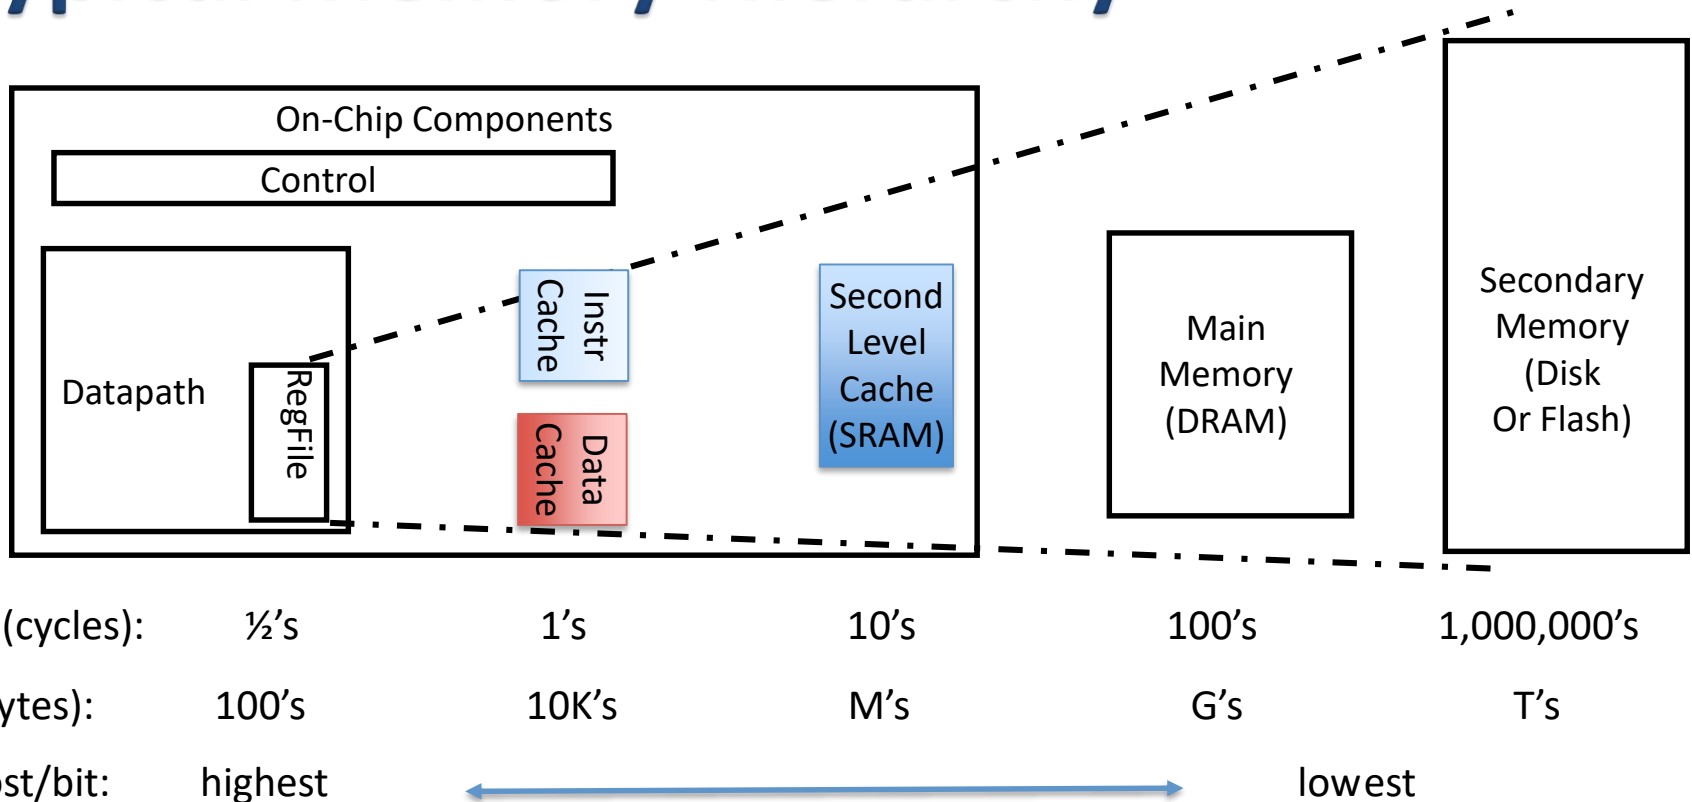
Announcement

- ▶ Lab #9
 - Due in 1 week
- ▶ Project #2
 - Start working on it during lab this week
 - Due Monday (4/29)
- ▶ HW #6 in CatCourses
 - Due Monday (4/22) at 11:59pm
- ▶ Reading assignment
 - Chapter 6.4-6.7 of zyBooks
 - Make sure to do the Participation Activities
 - Due Friday (4/19) at 11:59pm
 - Chapter 5.1-5.6 of zyBooks
 - Make sure to do the Participation Activities
 - Due Friday (4/26) at 11:59pm

Announcement

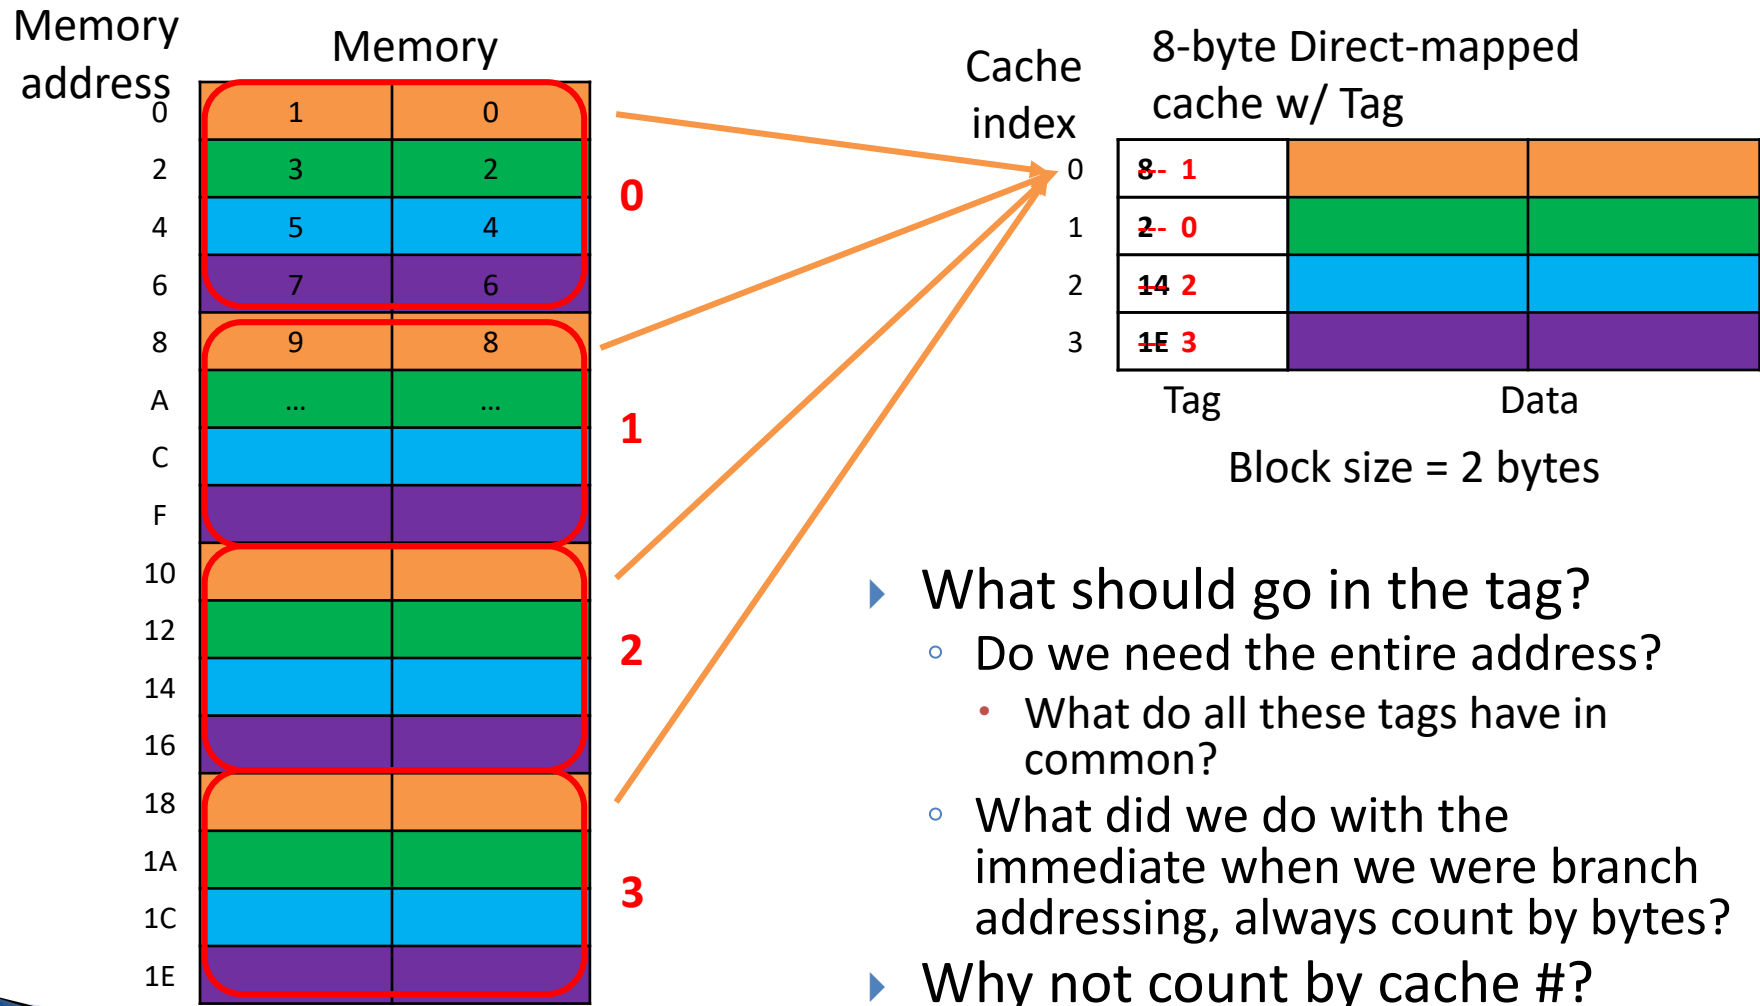
- ▶ Midterm Exam 2
 - 4/24 (Wednesday, in lecture) Not 4/17 as scheduled
 - Lectures #8 - #18
 - HW #2 - #6
 - Practice exam in CatCourses
 - Closed book
 - 1 sheet of note (8.5" x 11")
 - MIPS reference sheet will be provided

Typical Memory Hierarchy



- ▶ **Principle of locality + memory hierarchy** presents programmer with as much memory as is available in the *cheapest* technology at the speed offered by the *fastest* technology

Direct-Mapped Cache (4/4)



- ▶ What should go in the tag?
 - Do we need the entire address?
 - What do all these tags have in common?
 - What did we do with the immediate when we were branch addressing, always count by bytes?
- ▶ Why not count by cache #?
 - It's useful to draw memory with the same width as the block size

Issues with Direct-Mapped

- ▶ Since multiple memory addresses map to same cache index, how do we tell which one is in there?
- ▶ What if we have a block size > 1 byte?
 - Answer: divide memory address into three fields

t t t t t t t t t t t t t t t t	i i i i i i i i i i	o o o o
---------------------------------	---------------------	---------


Tag to check if it has the correct block

Index to select block

Byte
offset
within
block

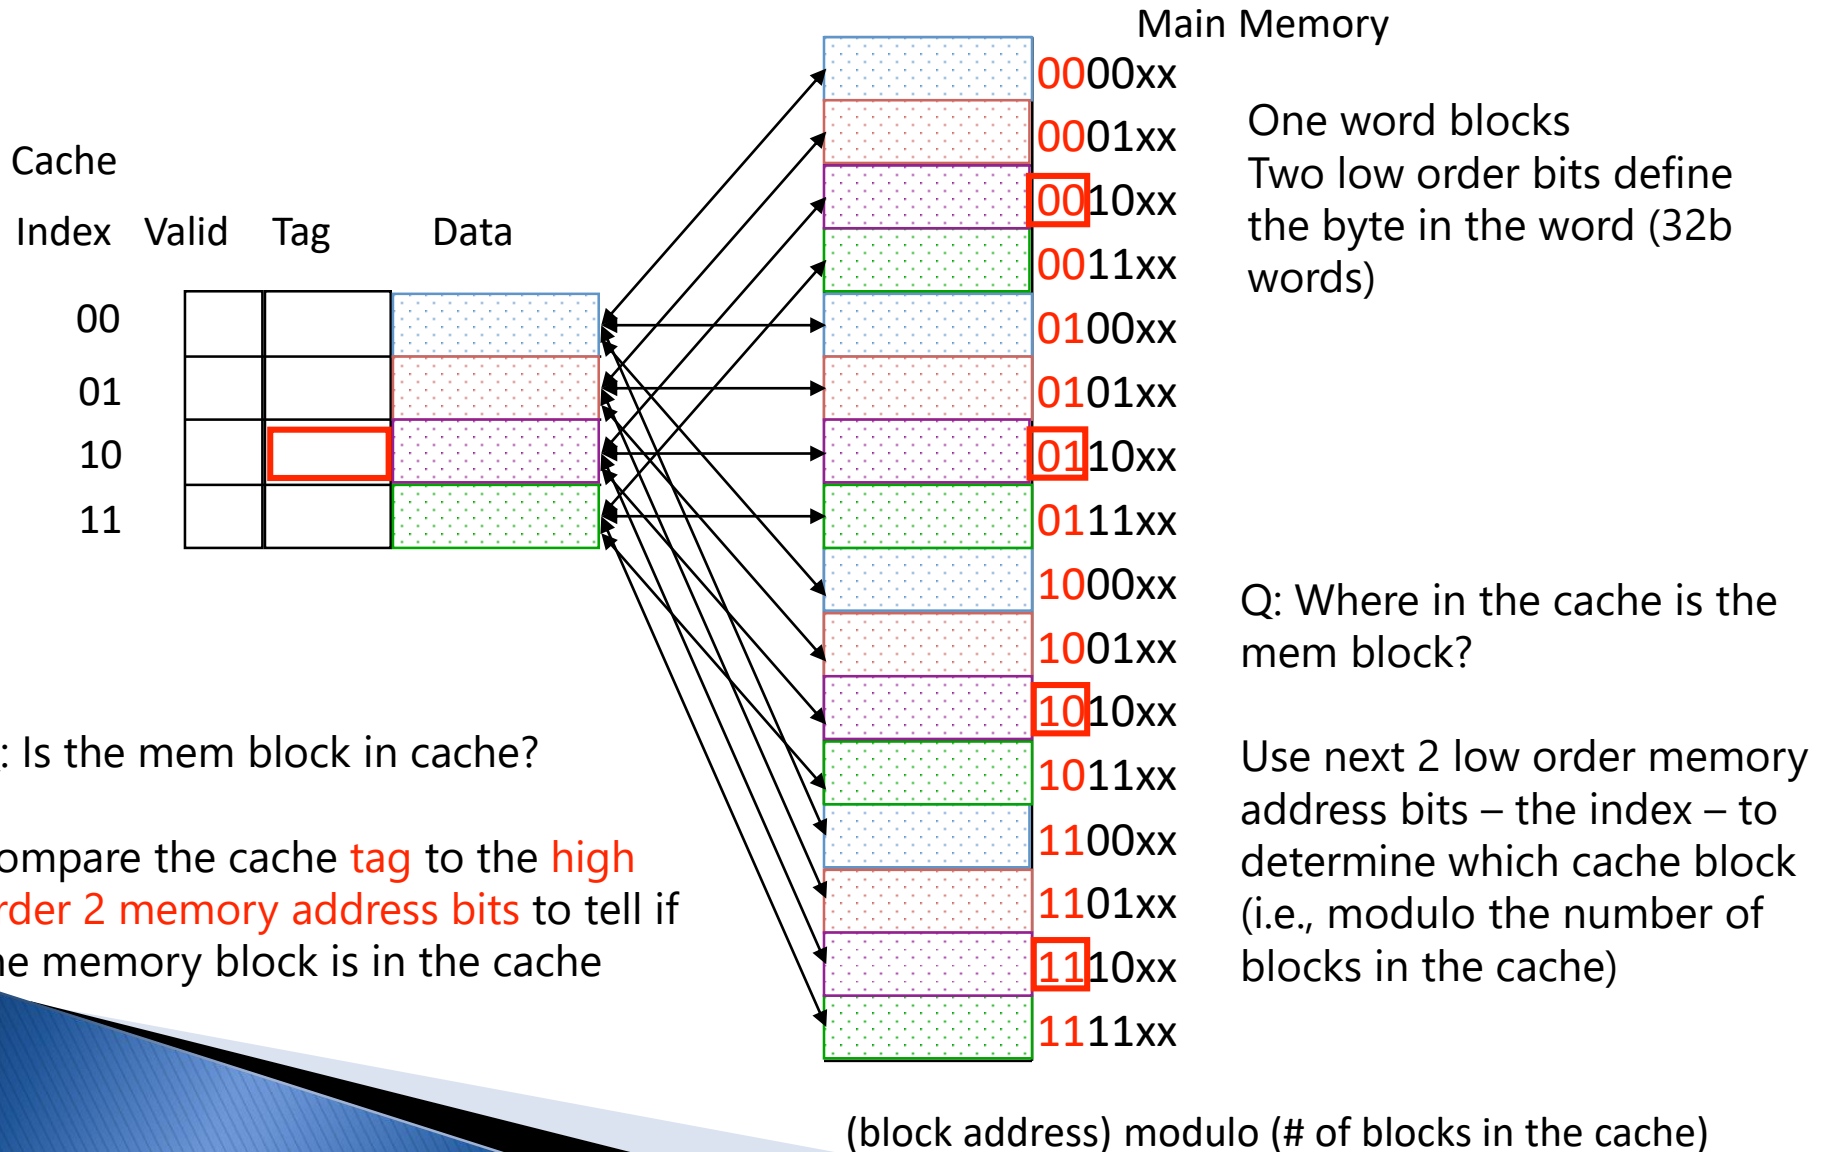
Direct Mapped Cache Example

One word blocks, cache size = 1K words (or 4KB)



What kind of locality are we taking advantage of?

Caching: A Simple First Example



Caching Terminology

- ▶ When reading memory, 3 things can happen:
 - **cache hit:**
 - cache block is valid and contains proper address, so read desired word
 - **cache miss:**
 - nothing in cache at appropriate block, so fetch from memory
 - **cache miss, block replacement:**
 - wrong data is in cache at appropriate block, so discard it and fetch desired data from memory (cache always copy)

Direct Mapped Cache

Consider the main memory word reference string

Start with an empty cache - all blocks
initially marked as not valid

0 1 2 3 4 3 4 15
0000 0001 0010 0011 0100 0011 0100 1111

Address 0

Direct Mapped Cache

Consider the main memory word reference string

Start with an empty 4-word cache - all blocks initially marked as not valid

0 1 2 3 4 3 4 15
0000 0001 0010 0011 0100 0011 0100 1111

0 miss

00	Mem(0)

- 1 requests, 1 miss

Direct Mapped Cache

Consider the main memory word reference string

Start with an empty cache - all blocks initially marked as not valid

0 1 2 3 4 3 4 15

0000 0001 0010 0011 0100 0011 0100 1111

0 miss

00	Mem(0)

1 miss

00	Mem(0)
00	Mem(1)

2 miss

00	Mem(0)
00	Mem(1)
00	Mem(2)

3 miss

00	Mem(0)
00	Mem(1)
00	Mem(2)
00	Mem(3)

4 miss

00	Mem(0)
00	Mem(1)
00	Mem(2)
00	Mem(3)

3 hit

01	Mem(4)
00	Mem(1)
00	Mem(2)
00	Mem(3)

4 hit

01	Mem(4)
00	Mem(1)
00	Mem(2)
00	Mem(3)

15 miss

01	Mem(4)
00	Mem(1)
00	Mem(2)
00	Mem(3)


11

15

- 8 requests, 6 misses

Multiword Block Direct Mapped Cache

Four words/block, cache size = 1K words



What kind of locality are we taking advantage of?

Taking Advantage of Spatial Locality

Let cache block hold more than one word

Start with an empty cache - all blocks initially marked as not valid

0 1 2 3 4 3 4 15

0 miss

00	Mem(1)	Mem(0)

1 hit

00	Mem(1)	Mem(0)

2 miss

00	Mem(1)	Mem(0)
00	Mem(3)	Mem(2)

3 hit

00	Mem(1)	Mem(0)
00	Mem(3)	Mem(2)

01 5 4 miss 4

00	Mem(1)	Mem(0)
00	Mem(3)	Mem(2)

3 hit

01	Mem(5)	Mem(4)
00	Mem(3)	Mem(2)

4 hit

01	Mem(5)	Mem(4)
00	Mem(3)	Mem(2)

15 miss


11

01	Mem(5)	Mem(4)
00	Mem(3)	Mem(2)

15 14

- 8 requests, 4 misses

Miss Rate vs Block Size vs Cache Size



Miss rate goes up if the block size becomes a significant fraction of the cache size because the number of blocks that can be held in the same size cache is smaller (increasing capacity misses)

Average Memory Access Time (AMAT)

- ▶ Average Memory Access Time (AMAT) is the average to access memory considering both hits and misses

$$\text{AMAT} = \text{Time for a hit} + \text{Miss rate} \times \text{Miss penalty}$$

- ▶ What is the AMAT for a processor with a 200 psec clock, a miss penalty of 50 clock cycles, a miss rate of 0.02 misses per instruction and a cache access time of 1 clock cycle?

$$1 + 0.02 \times 50 = 2 \text{ clock cycles}$$

$$\text{Or } 2 \times 200 = 400 \text{ psecs}$$

- ▶ Potential impact of much larger cache on AMAT?

- 1) Lower Miss rate
- 2) Longer Access time (Hit time): smaller is faster

At some point, increase in hit time for a larger cache may overcome the improvement in hit rate, yielding a decrease in performance

Block Size Tradeoff (1/3)

- ▶ Benefits of Larger Block Size
 - **Spatial Locality:** if we access a given word, we're likely to access other nearby words soon
 - Very applicable with Stored-Program Concept: if we execute a given instruction, it's likely that we'll execute the next few as well
 - Works nicely in sequential array accesses too

Block Size Tradeoff (2/3)

▶ Drawbacks of Larger Block Size

- Larger block size means **larger miss penalty**
 - on a miss, takes longer time to load a new block from next level
- If block size is too big relative to cache size, then there are too few blocks
 - Result: miss rate goes up

▶ In general, minimize

Average Memory Access Time (AMAT)

= Hit Time + Miss Penalty x Miss Rate

Block Size Tradeoff (3/3)

- ▶ Hit Time
 - time to find and retrieve data from current level cache
- ▶ Miss Penalty
 - average time to retrieve data on a current level miss (includes the possibility of misses on successive levels of memory hierarchy)
- ▶ Hit Rate
 - % of requests that are found in current level cache
- ▶ Miss Rate
 - $1 - \text{Hit Rate}$

Extreme Example: One Big Block

Valid Bit




Tag

Cache Data



- ▶ Cache Size = 4 bytes Block Size = 4 bytes
 - Only **ONE** entry (row) in the cache!
- ▶ If item accessed, likely accessed again soon
 - But unlikely will be accessed again immediately!
- ▶ The next access will likely to be a miss again
 - Continually loading data into the cache but discard data (force out) before use it again
 - Nightmare for cache designer: **Ping Pong Effect**

Block Size Tradeoff Conclusions



What to do on a write hit?

▶ Write-through

- update the word in cache block and corresponding word in memory

▶ Write-back

- update word in cache block
- allow memory word to be “stale”
- add ‘dirty’ bit to each block indicating that memory needs to be updated when block is replaced
- OS flushes cache before I/O...

▶ Performance trade-offs?

Types of Cache Misses (1/2)

- ▶ “Three Cs” Model of Misses
- ▶ 1st C: **Compulsory Misses**
 - occur when a program is first started
 - cache does not contain any of that program’s data yet, so misses are bound to occur (valid bit = 0)
 - can’t be avoided easily, so won’t focus on these in this course

Types of Cache Misses (2/2)

▶ 2nd C: Conflict Misses

- miss that occurs because two distinct memory addresses map to the same cache location
- two blocks (which happen to map to the same location) can keep overwriting each other
- big problem in direct-mapped caches
- how do we lessen the effect of these?

▶ Dealing with Conflict Misses

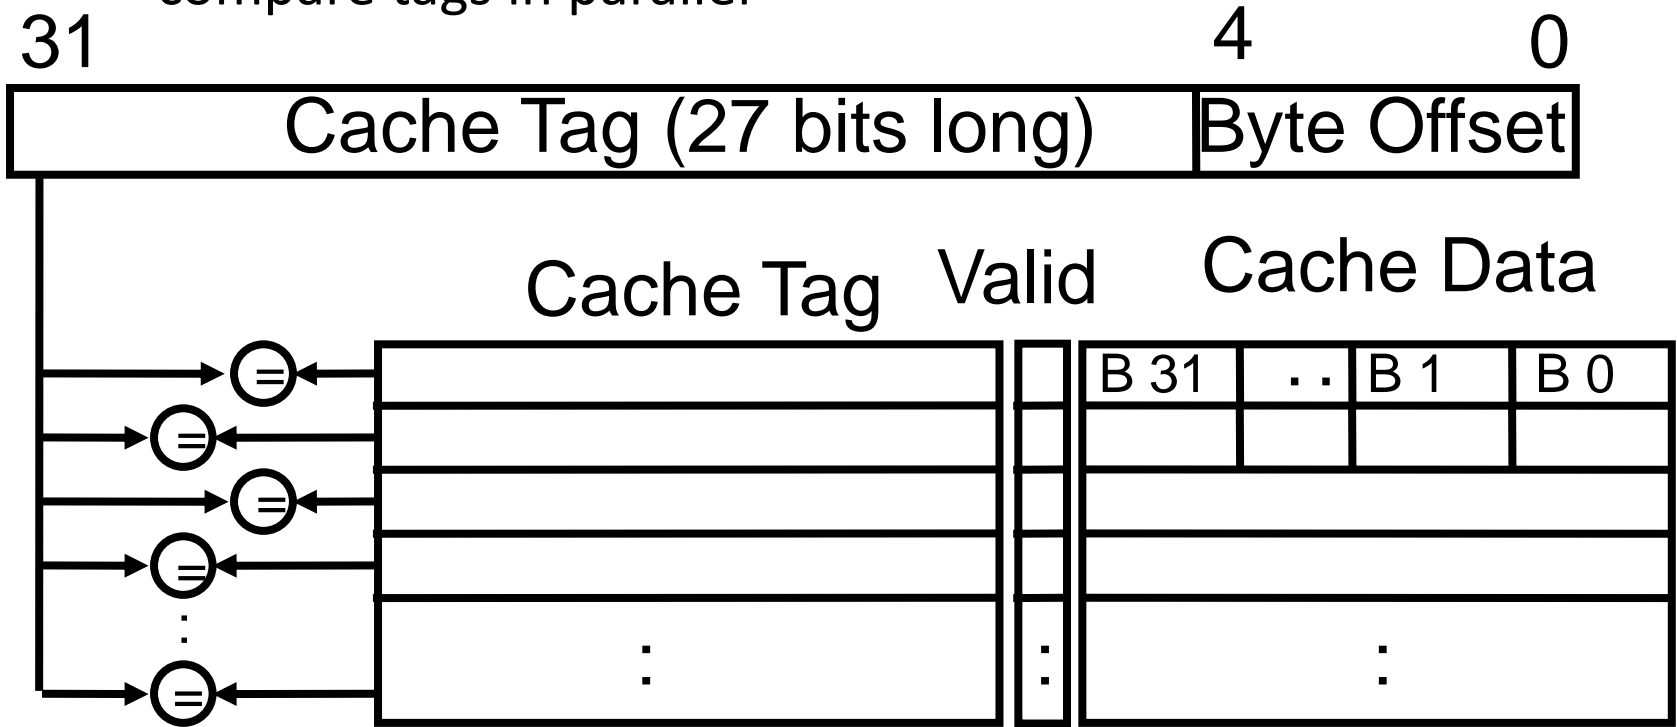
- Solution 1: Make the cache size bigger
 - Fails at some point
- Solution 2: Multiple distinct blocks can fit in the same cache Index
 - How???

Fully Associative Cache (1/3)

- ▶ Memory address fields:
 - **Tag**: same as before
 - **Offset**: same as before
 - **Index**: non-existent
- ▶ What does this mean?
 - no “rows”: any block can go anywhere in the cache
 - must compare with all tags in entire cache to see if data is there

Fully Associative Cache (2/3)

- ▶ Fully Associative Cache (e.g., 32 B block)
 - compare tags in parallel



Fully Associative Cache (3/3)

- ▶ Benefit of Fully Assoc. Cache
 - No Conflict Misses (since data can go anywhere)
- ▶ Drawbacks of Fully Assoc. Cache
 - Need hardware comparator for every single entry: if we have a 64KB of data in cache with 4B entries, we need 16K comparators: infeasibly high cost

Final Type of Cache Miss

- ▶ 3rd C: Capacity Misses
 - miss that occurs because the cache has a limited size
 - miss that would not occur if we increase the size of the cache
 - sketchy definition, so just get the general idea
- ▶ This is the primary type of miss for Fully Associative caches.

N-Way Set Associative Cache (1/3)

- ▶ Memory address fields:
 - **Tag**: same as before
 - **Offset**: same as before
 - **Index**: points us to the correct “row” (called a set in this case)
- ▶ So what’s the difference?
 - each set contains multiple blocks
 - once we’ve found correct set, must compare with all tags in that set to find our data
 - Hybrid of direct-mapped and fully associative