

CSE 140

Computer Architecture

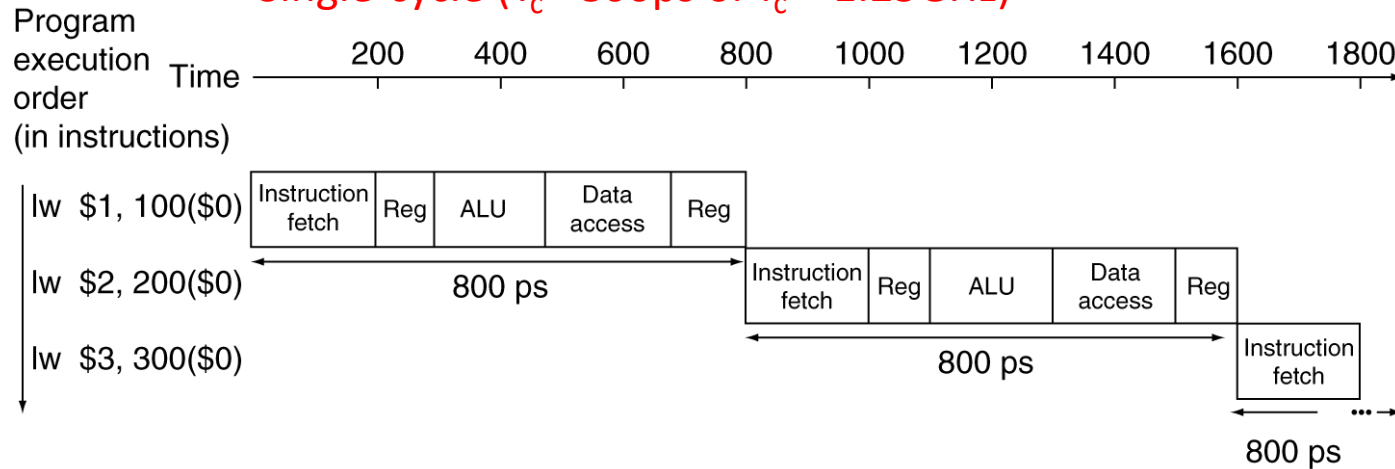
Lecture 5 – Pipelining (2)

Announcement

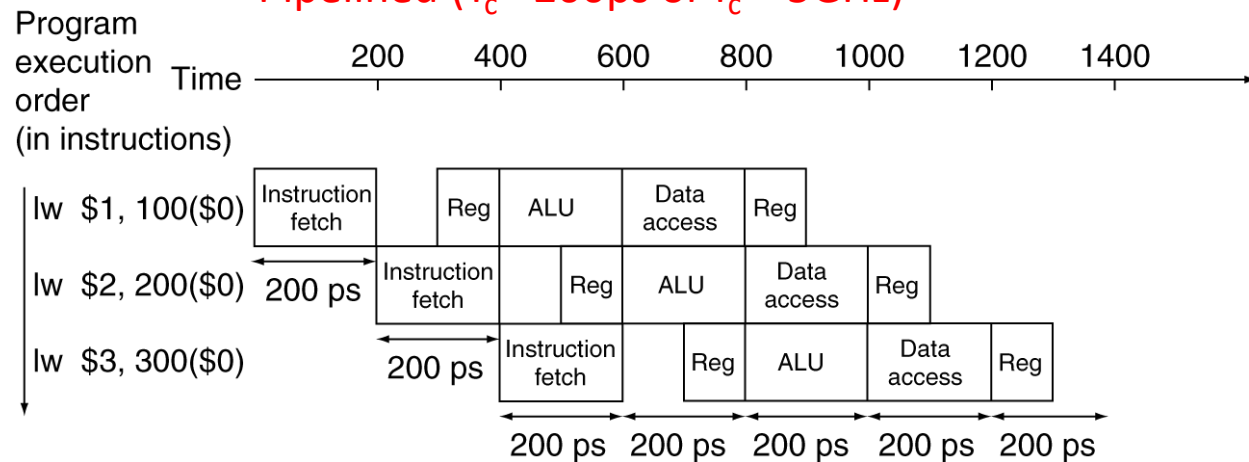
- ▶ Lab #1 this week
 - Due in one week
- ▶ Project #1 out this Friday
 - Due 10/11 (Monday) at 11:59pm
- ▶ Reading assignment #3
 - Chapter 5.1 – 5.4
 - Do all **Participation Activities** in each section
 - Access through CatCourses
 - Due Tuesday (9/19) at 11:59pm

Pipeline Performance

Single-cycle ($T_c = 800\text{ps}$ or $f_c = 1.25\text{GHz}$)



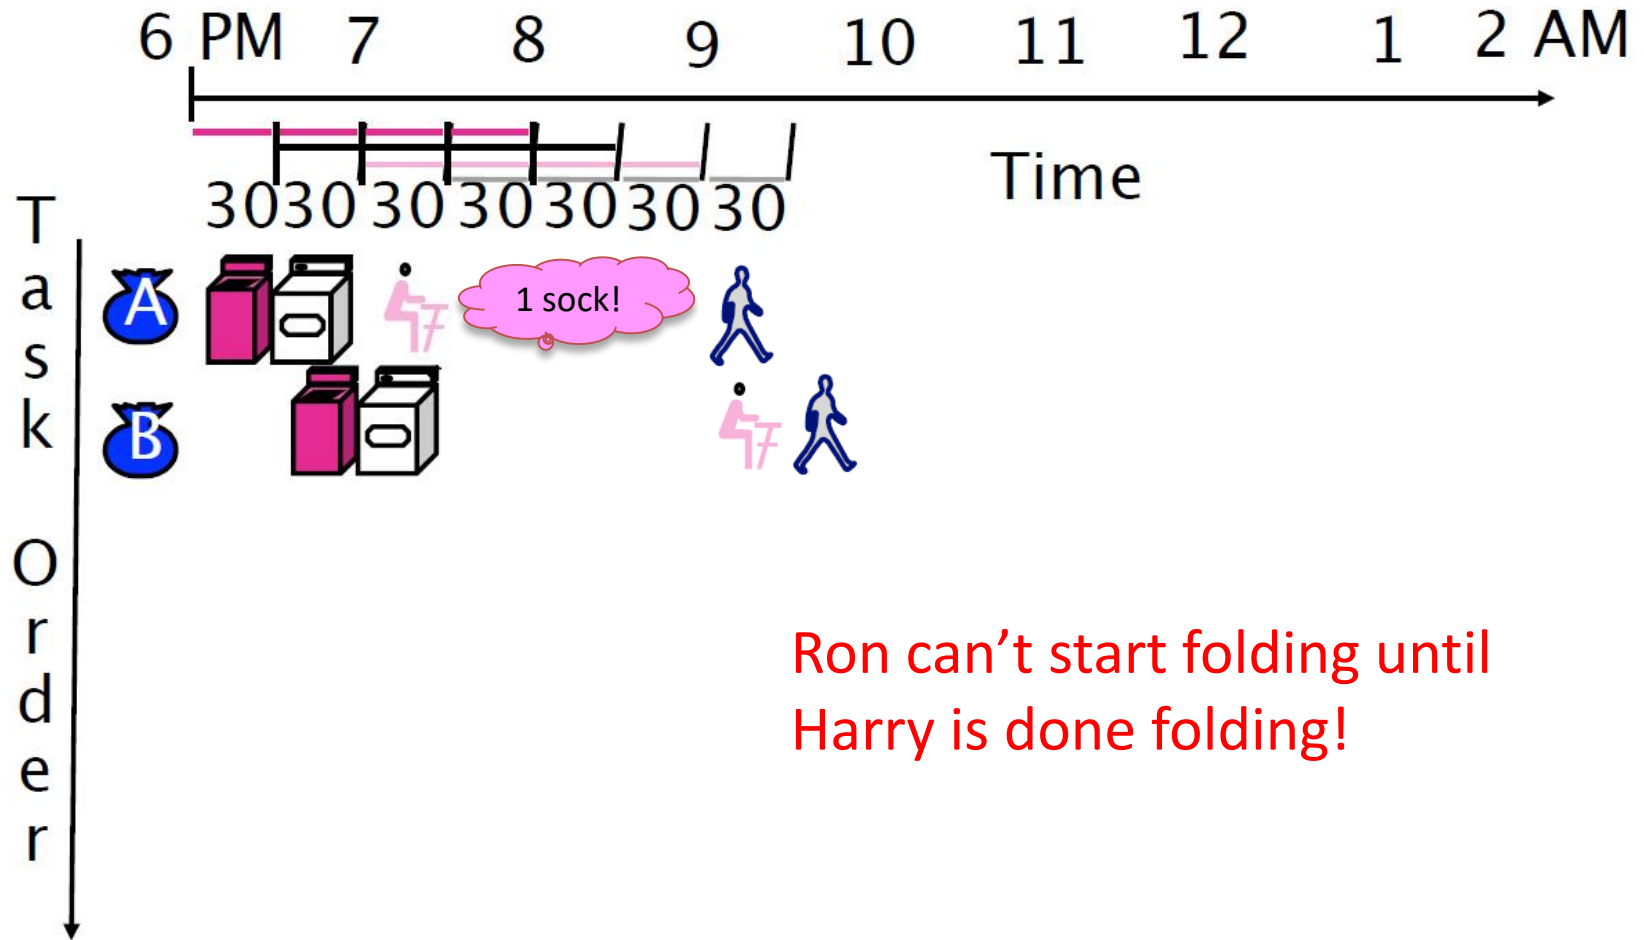
Pipelined ($T_c = 200\text{ps}$ or $f_c = 5\text{GHz}$)



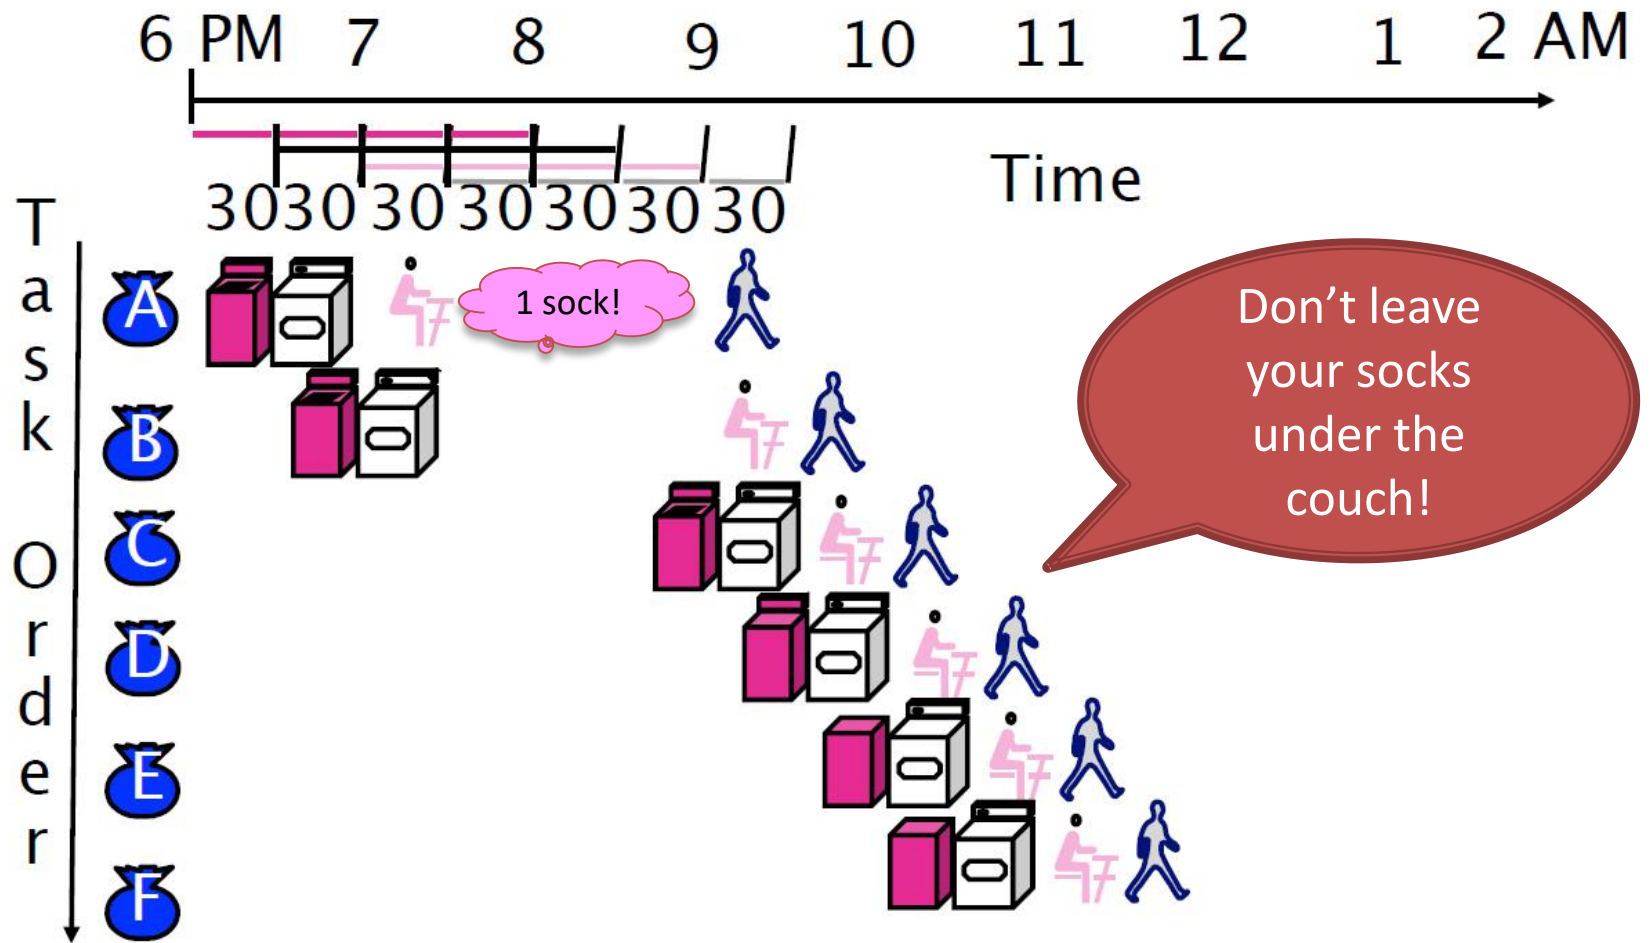
Pipeline Speedup

- ▶ If all stages are balanced
 - i.e., all take the same time
 - Time between instructions_{pipelined}
=
$$\frac{\text{Time between instructions}_{\text{nonpipelined}}}{\text{Number of stages}}$$
- ▶ If not balanced, speedup is less
- ▶ Speedup due to increased throughput
 - Latency (time for each instruction) does not decrease

Pipeline Hazard:



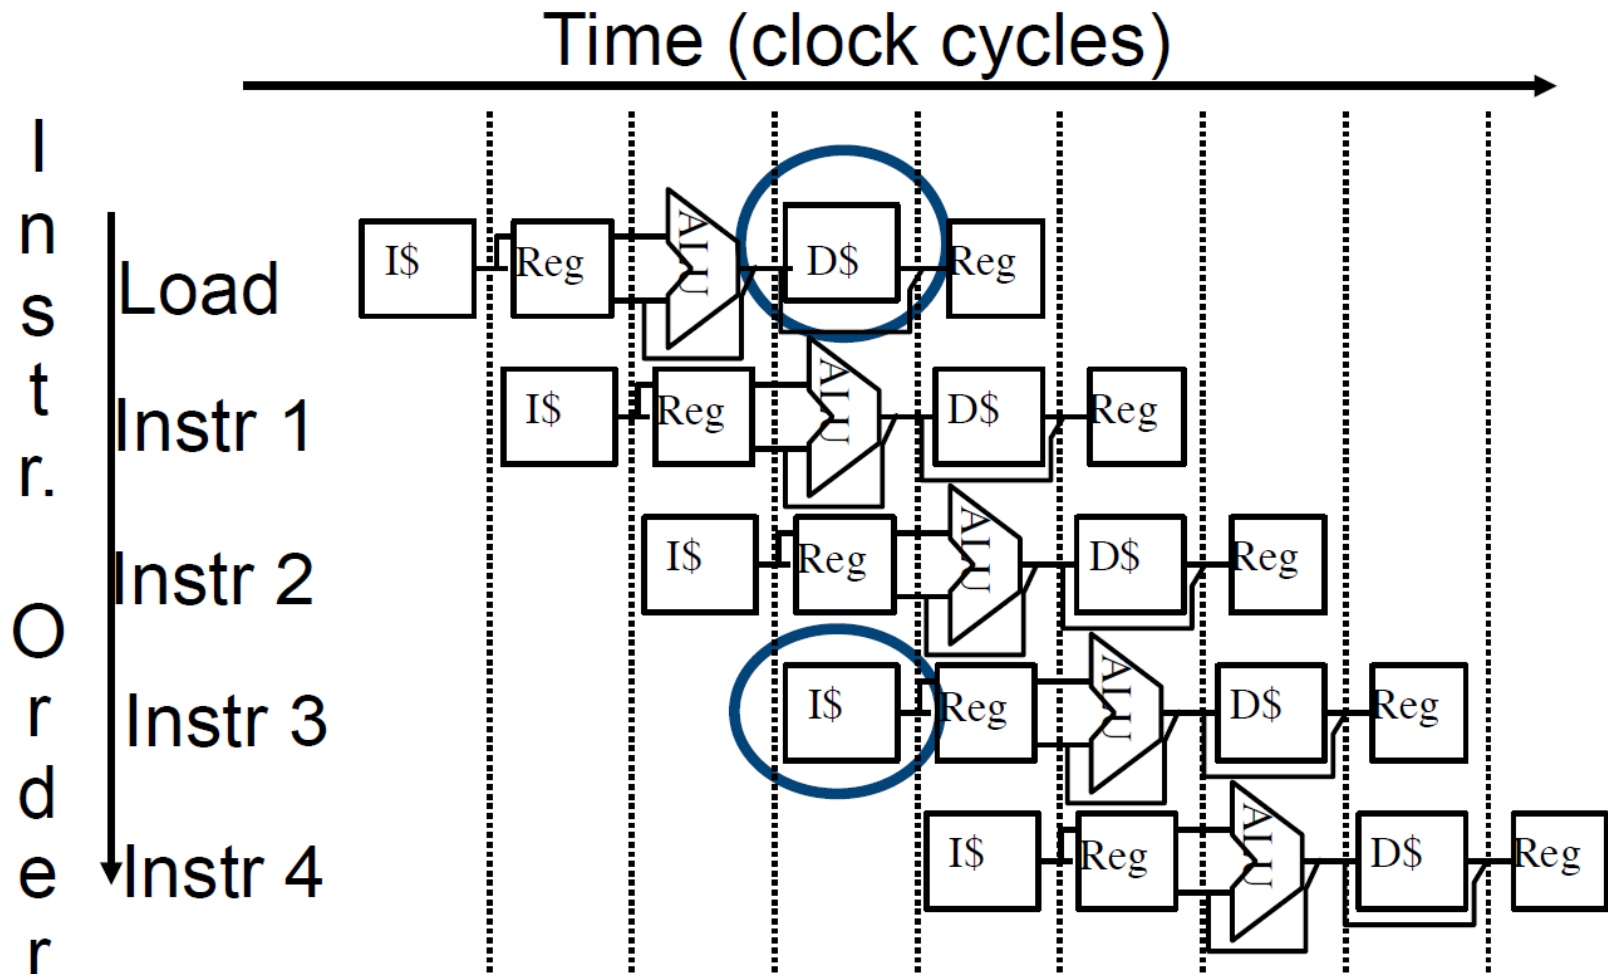
Pipeline Hazard:



Problems for Pipelining CPUs

- ▶ Limits to pipelining: Hazards prevent next instruction from executing during its designated clock cycle
 - **Structural hazards**: HW cannot support some combination of instructions (There is only one table for folding and stashing)
 - **Data hazards**: Instruction depends on result of prior instruction still in the pipeline (missing sock)
 - **Control hazards**: Pipelining of branches causes later instruction fetches to wait for the result of the branch
- ▶ These might result in pipeline **stalls** or “**bubbles**” in the pipeline.

Structural Hazard #1: Single Memory



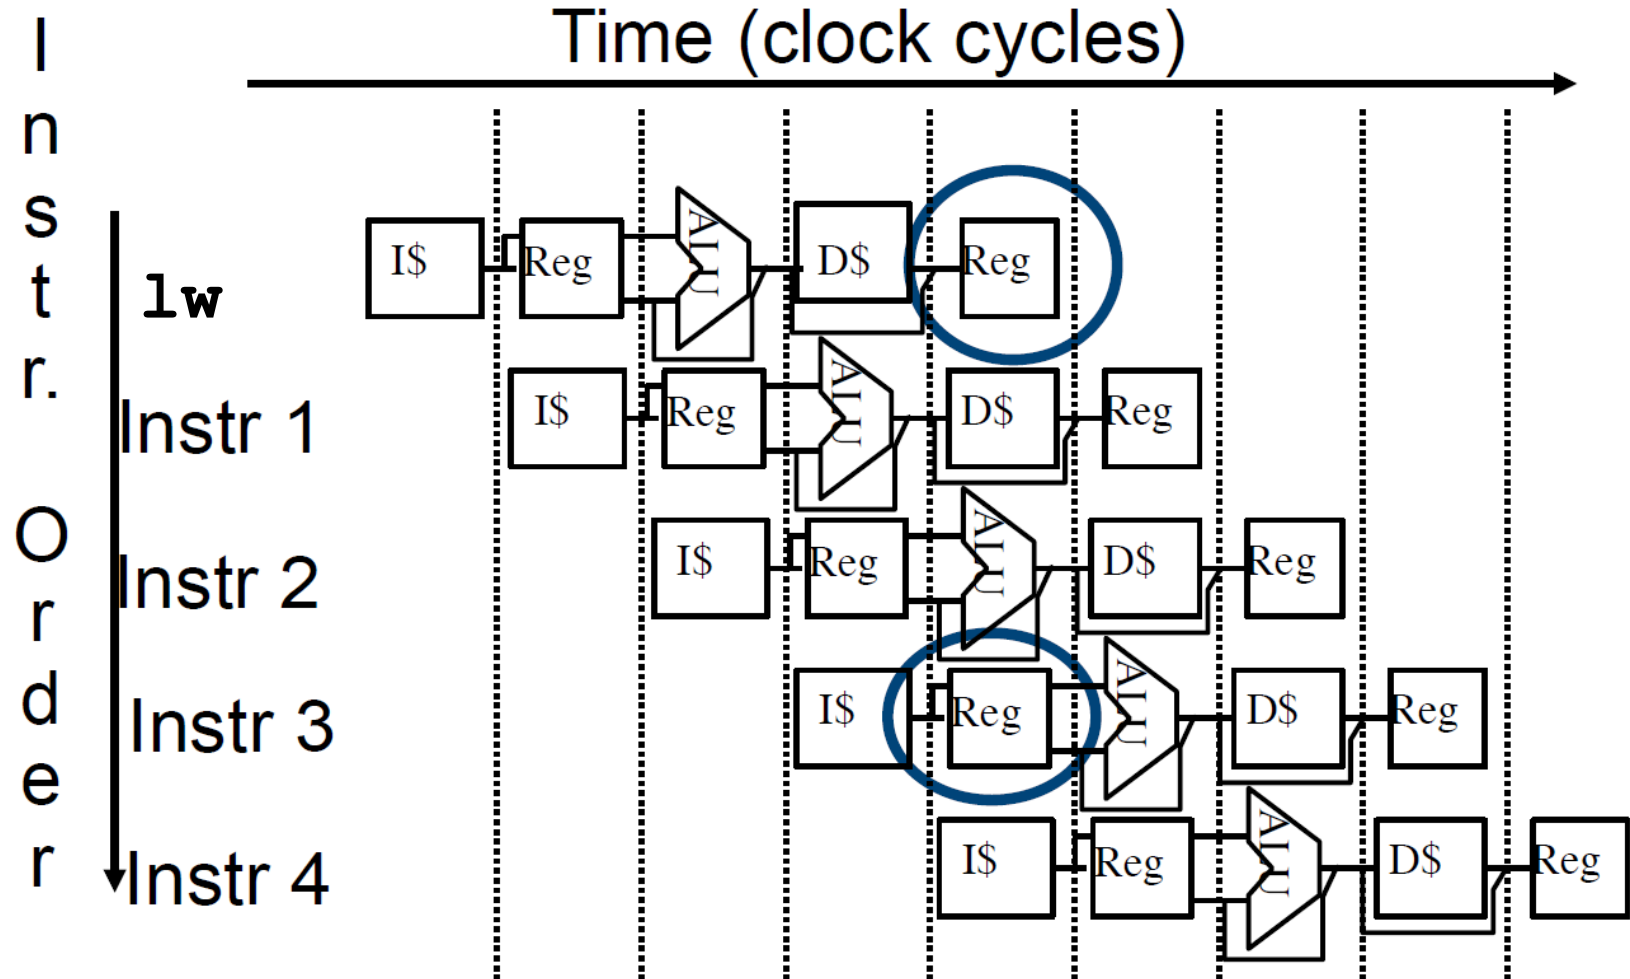
- ▶ Read same memory twice in the same cycle

Structural Hazard #1: Single Memory

▶ Solution:

- Infeasible and inefficient to create second memory
 - (We'll learn more about this next week)
- Simulate this by having two Level 1 Caches (a temporary smaller [of usually most recently used] copy of memory)
 - Have both an L1 Instruction Cache and an L1 Data Cache
 - Need more complex hardware to control when both caches miss

Structural Hazard #2: Registers



▶ Can we read and write to registers simultaneously?

Structural Hazard #2: Registers

- ▶ Two different solutions have been used:
 1. RegFile access is VERY fast: takes less than half the time of ALU stage
 - Write to Registers during first half of each clock cycle
 - Read from Registers during second half of each clock cycle
 2. Build RegFile with independent read and write ports
- ▶ Result: can perform Read and Write during same clock cycle

Data Hazards

- ▶ Consider this sequence:

sub \$2, \$1, \$3

and \$12, \$2, \$5

or \$13, \$6, \$2

add \$14, \$2, \$2

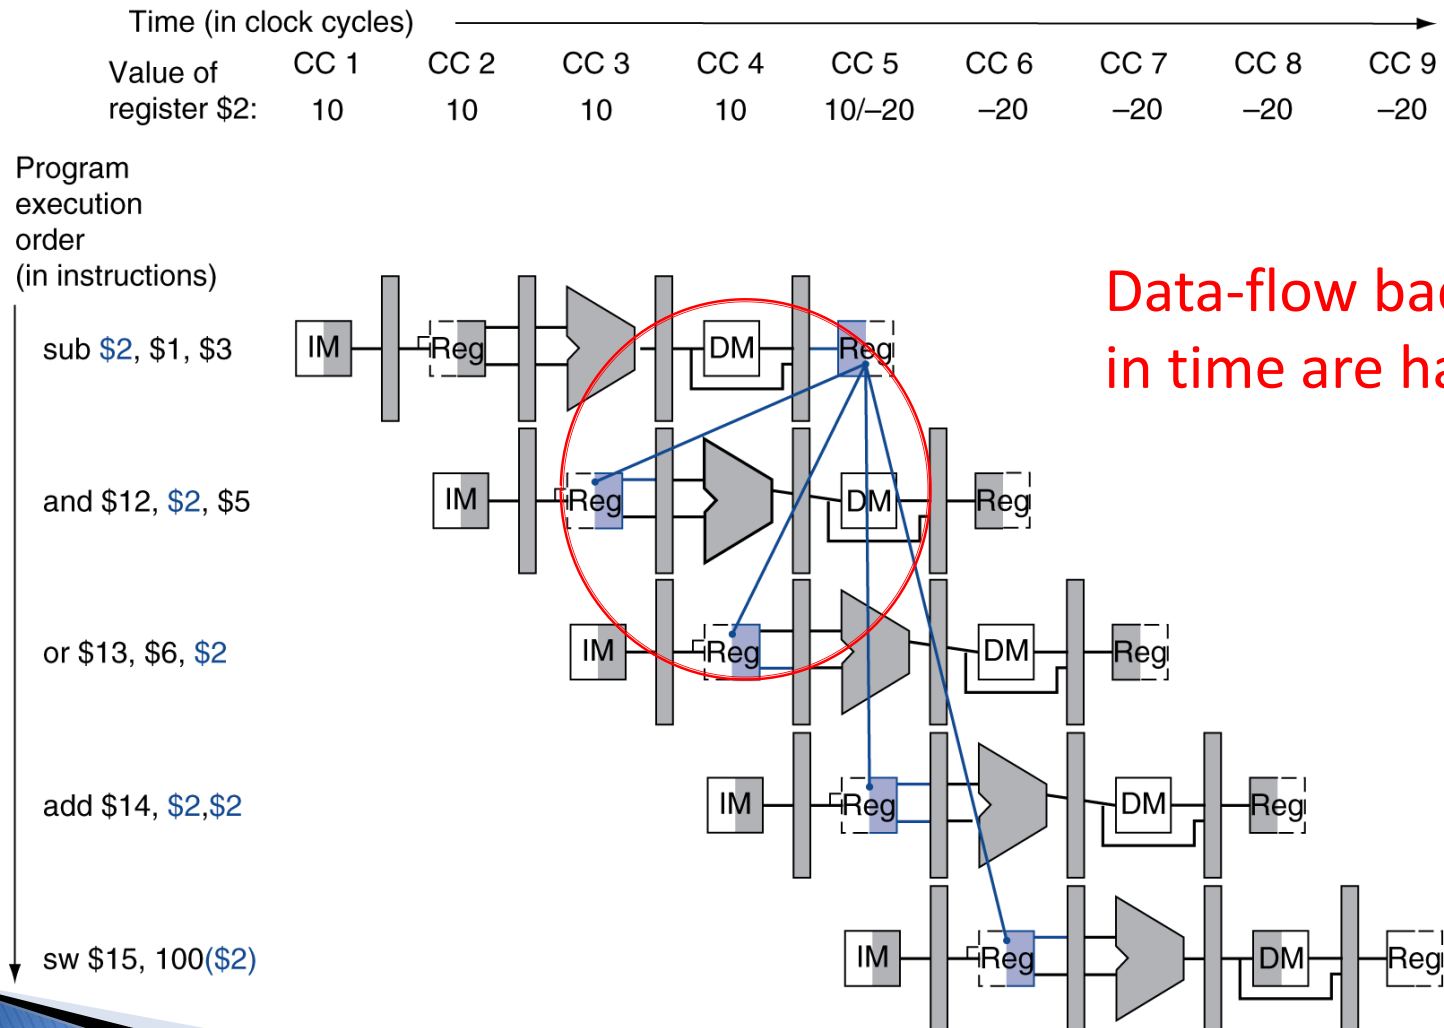
sw \$15, 100(\$2)



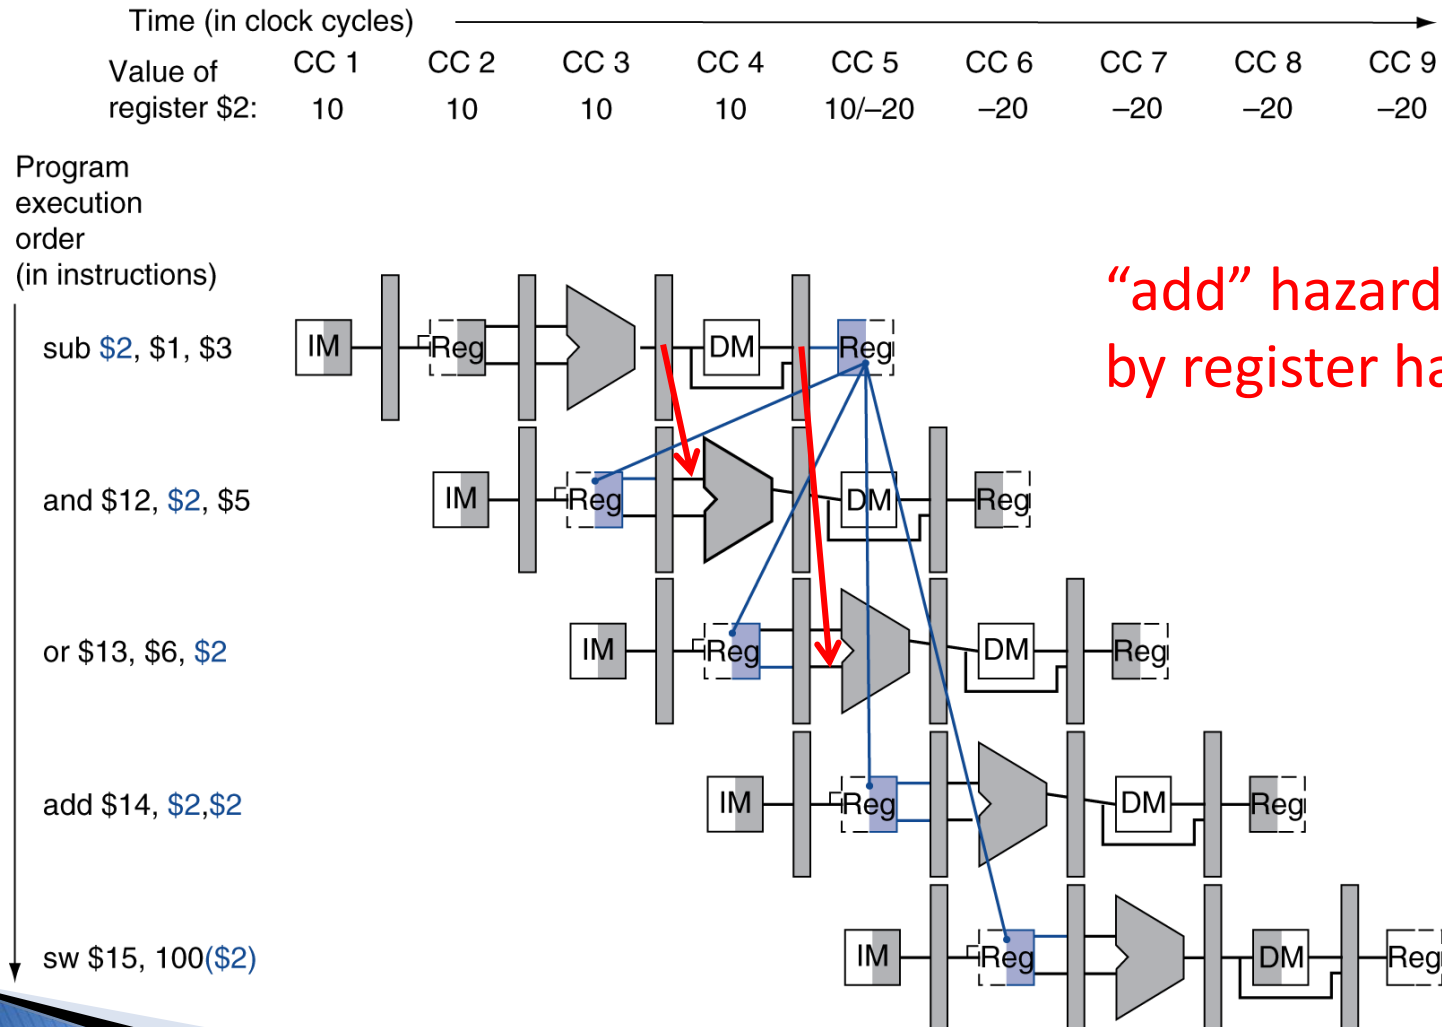
Can we read \$2 during ID stage?

What stage is the last instr. in at this moment?

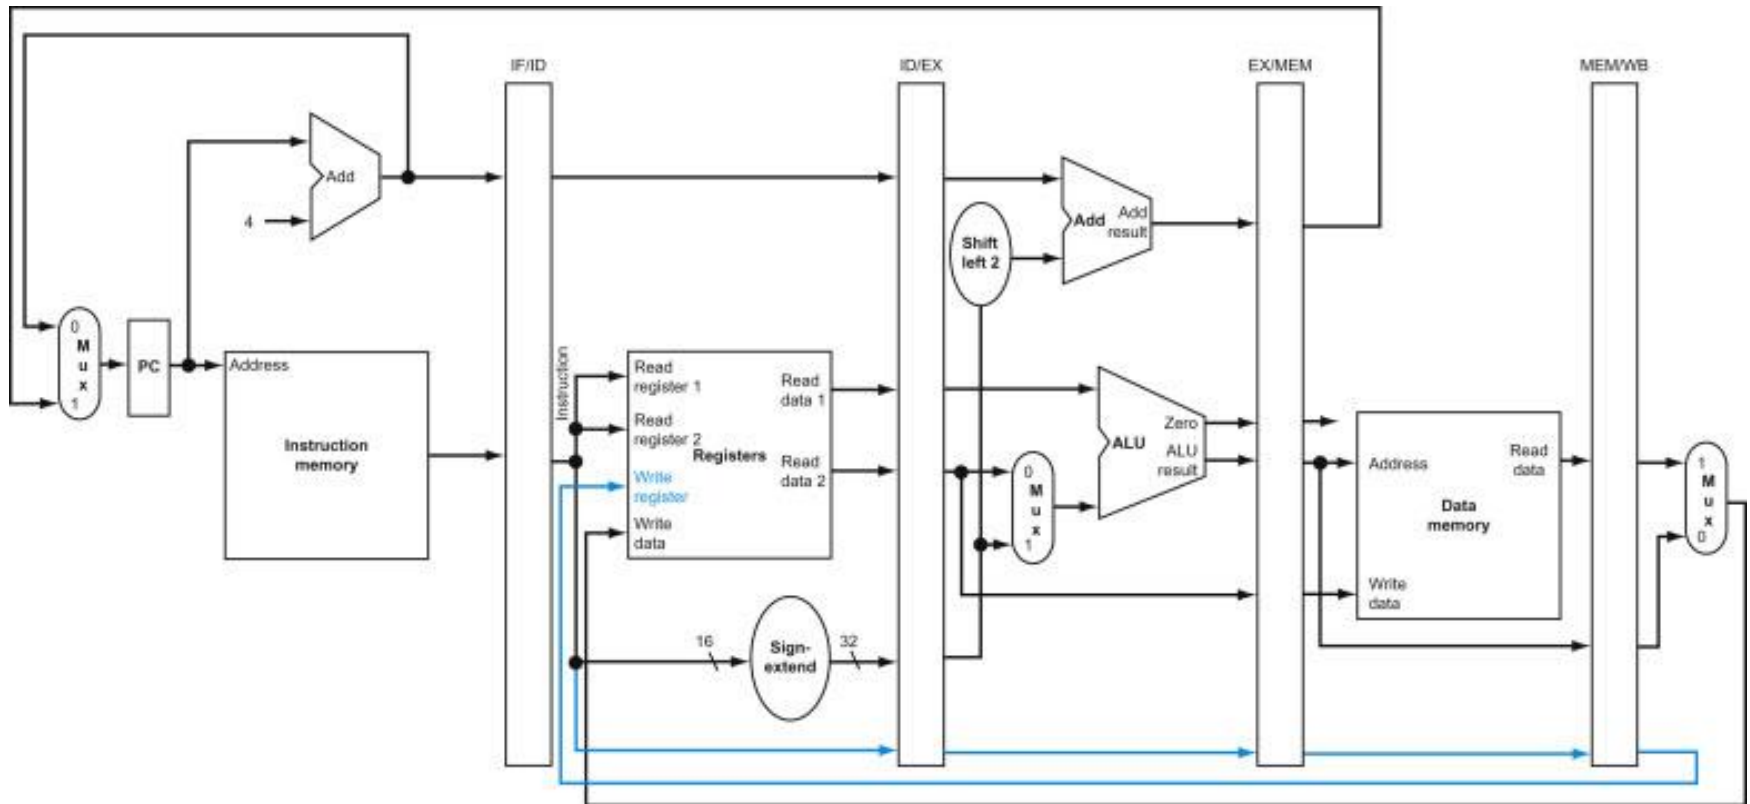
Data Hazards



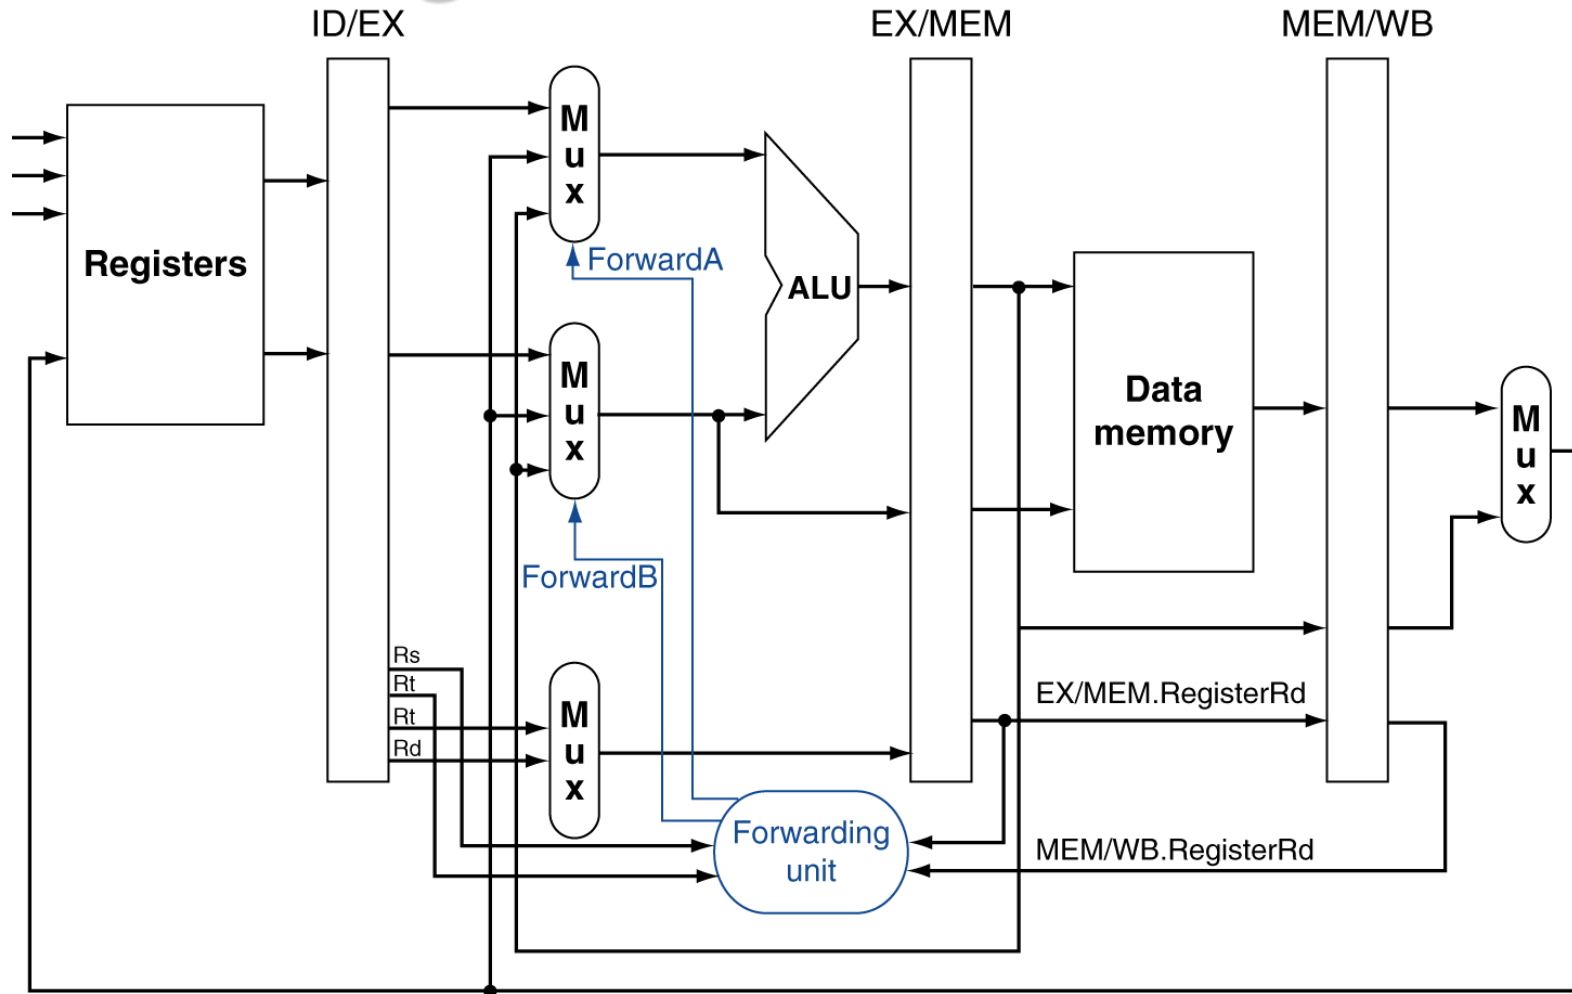
Data Hazard Solution: Forwarding



How do we Forward Data?

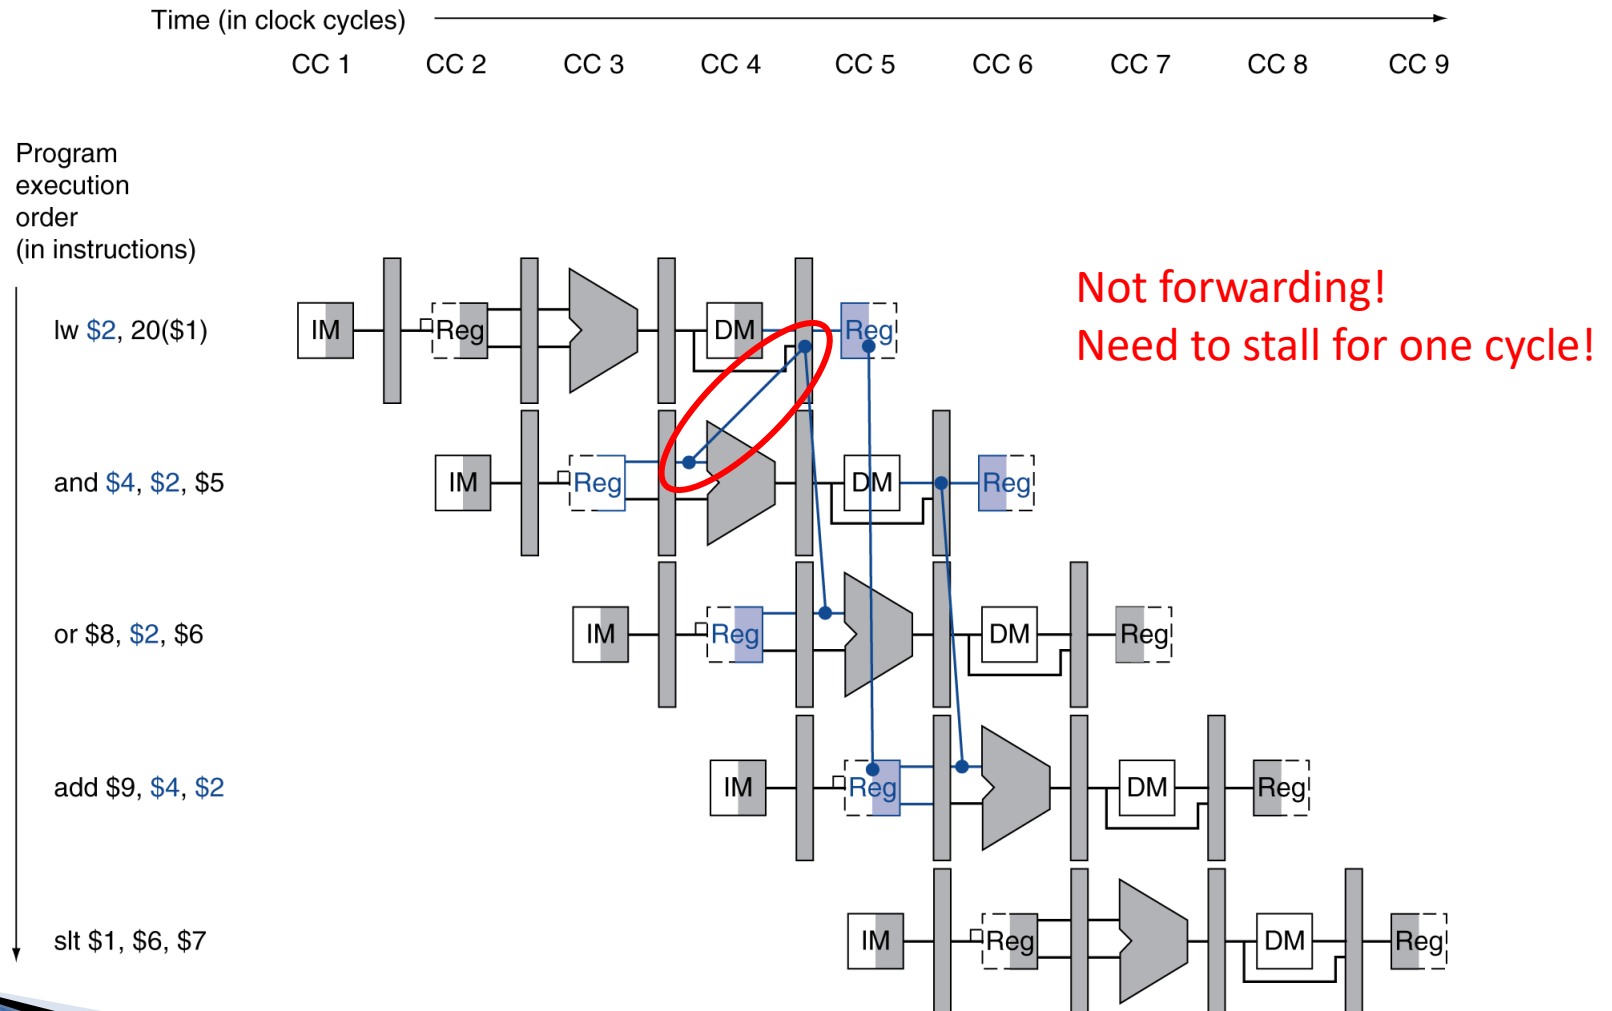


Forwarding Paths



b. With forwarding

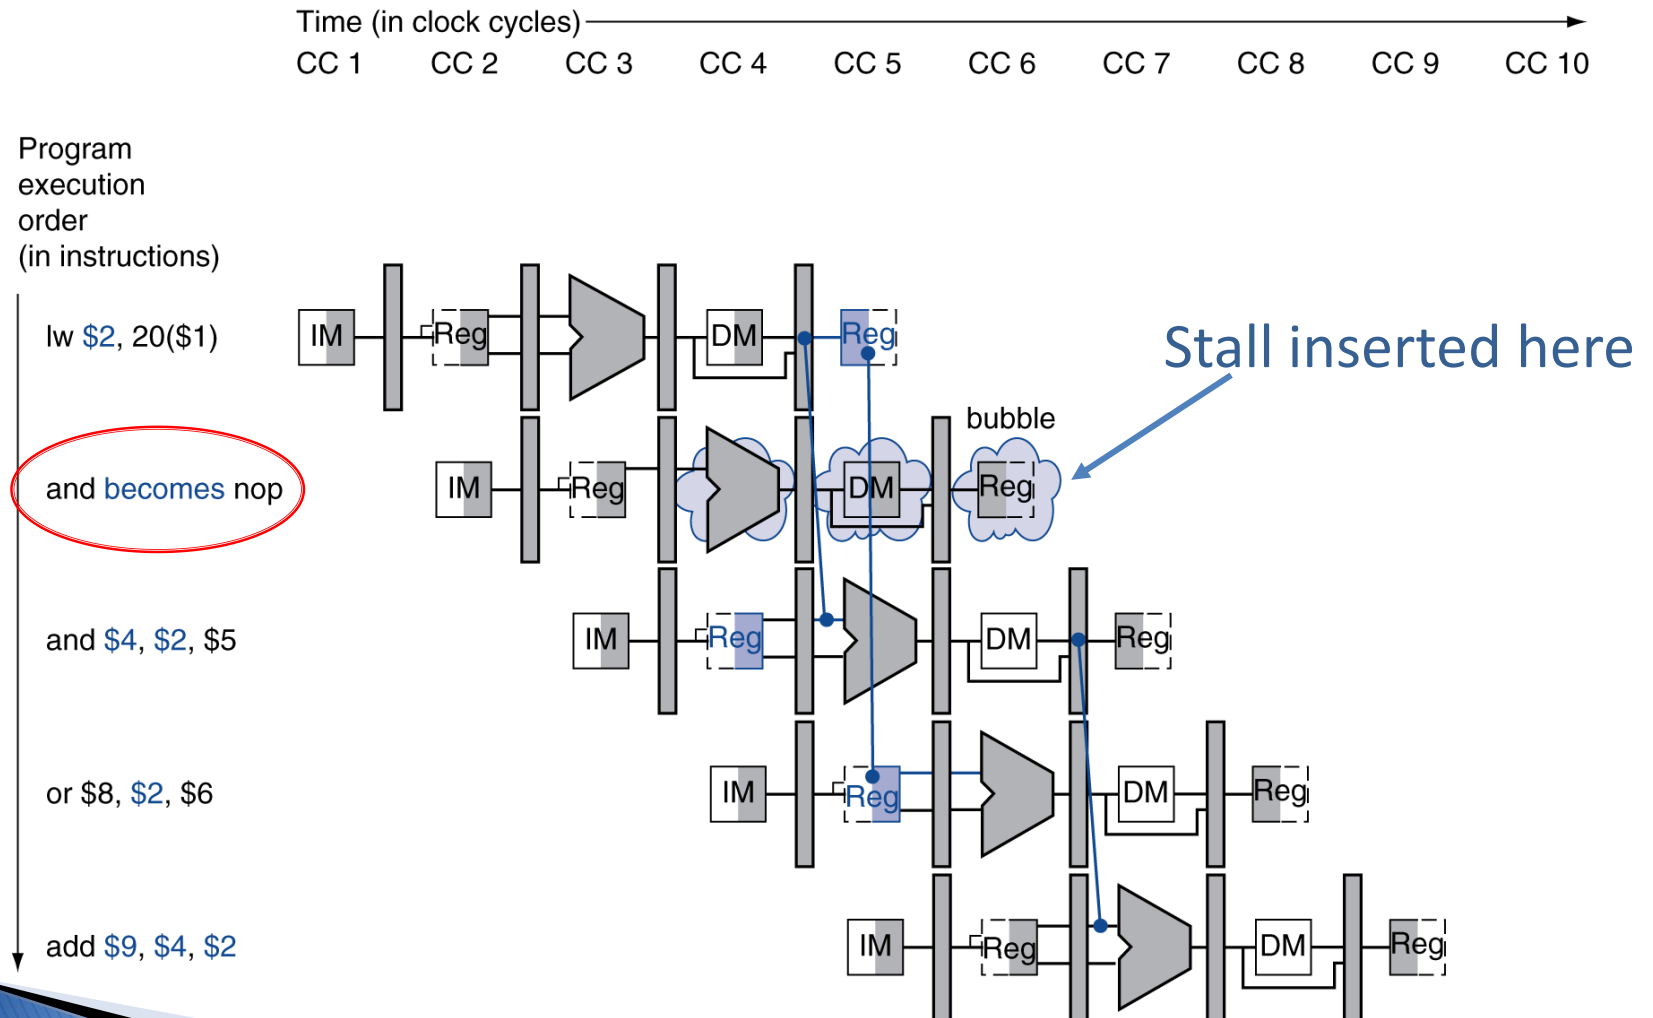
Data Hazard: Loads



Data Hazard: Loads

- ▶ Instruction slot after a load is called “load delay slot”
- ▶ If an instruction uses the result of the load, then the hardware interlock will stall it for one cycle.
- ▶ If the compiler puts an unrelated instruction in that slot, then no stall
- ▶ Letting the hardware stall the instruction in the delay slot is equivalent to putting a **nop** in the slot (except the latter uses more code space)

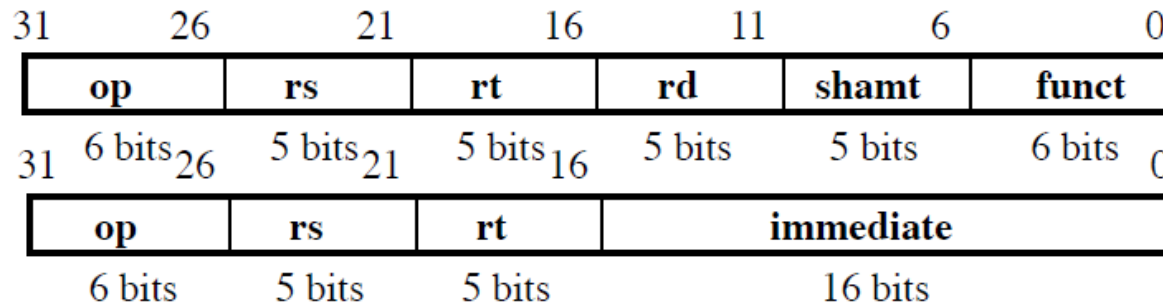
Data Hazard: Loads



Pipelining and ISA Design

- ▶ MIPS ISA (Instruction Set Architecture) designed for pipelining
 - All instructions are 32-bits
 - Easier to fetch and decode in one cycle
 - x86: 1- to 17-byte instructions
 - HW translates instructions to internal RISC instructions
 - Few and regular instruction formats
 - Can decode and read registers in one step
 - Load/store addressing
 - Can calculate address in 3rd stage, access memory in 4th stage
 - Alignment of memory operands
 - Memory access takes only one cycle


Why Isn't the Destination Register Always in the Same Field in MIPS ISA?




- ▶ Need to have 2 part immediate if 2 sources and 1 destination always in same place!

SPUR processor
(Alternative ISA)


Register-Register: Rd, Rs1, Rs2



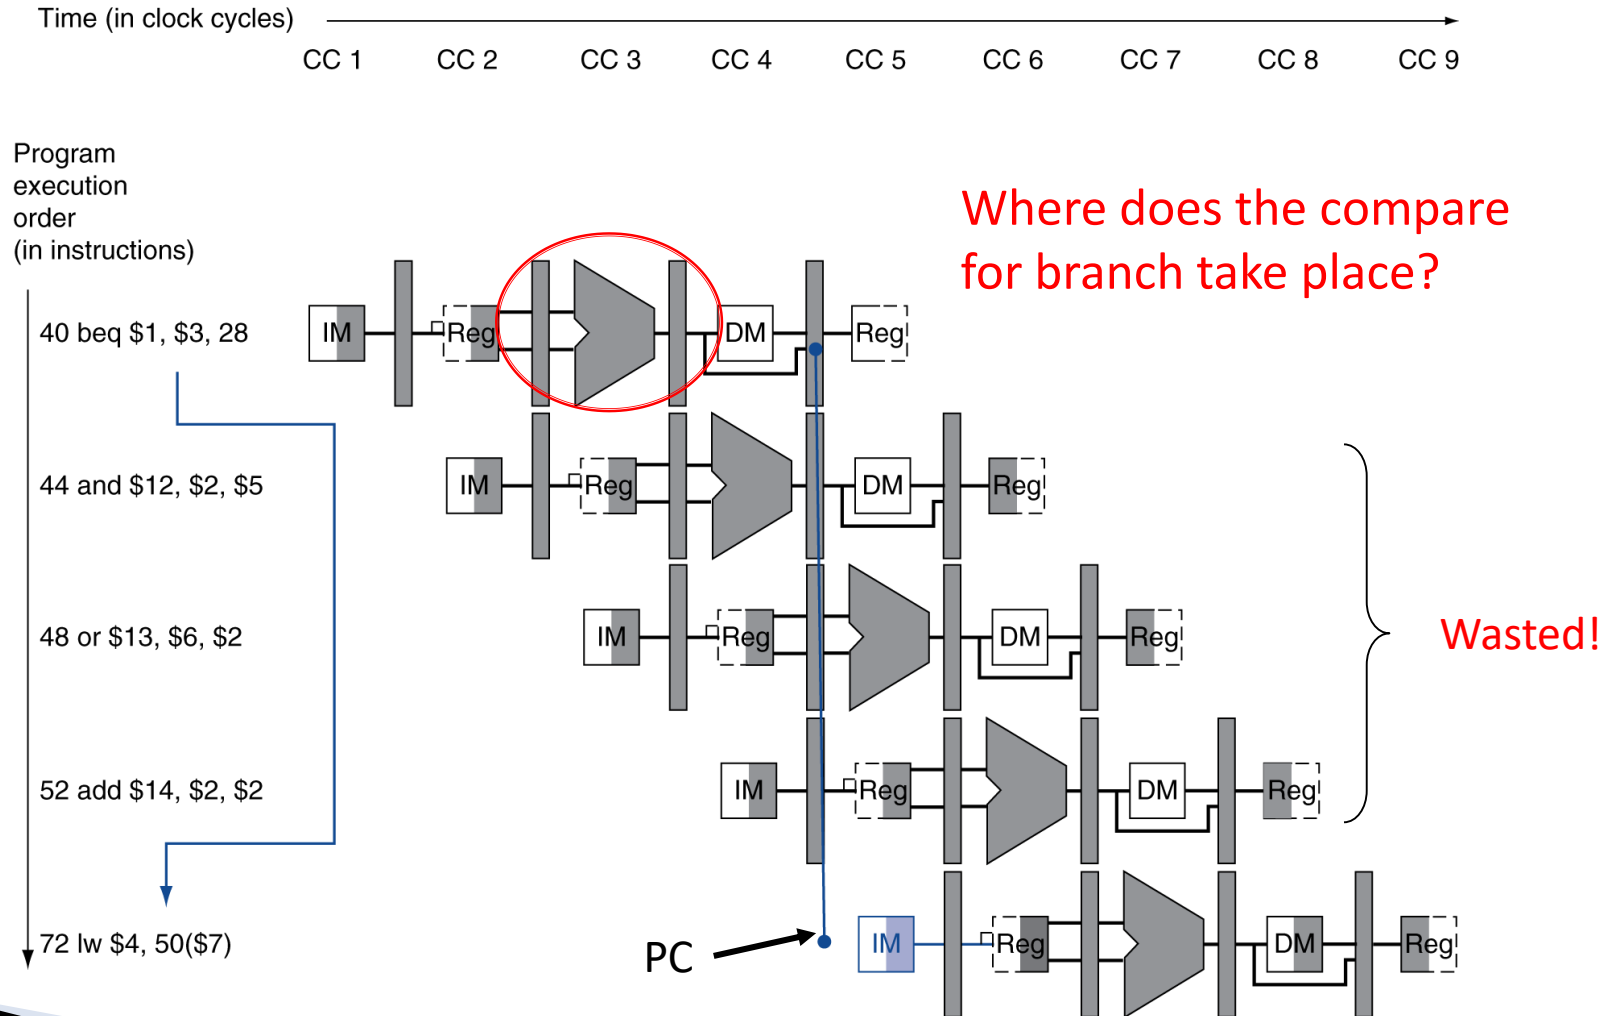
Register-Register: Rd, Rs1, Immediate



Store: Rs2, Rs1, Immediate



Control Hazard: Branching



Control Hazard: Branching

- ▶ We had put branch decision-making hardware in ALU stage
 - Therefore two more instructions after the branch will always be fetched, whether or not the branch is taken
- ▶ Desired functionality of a branch (ideal case)
 - If we do not take the branch, don't waste any time and continue executing normally
 - If we take the branch, don't execute any instructions after the branch, just go to the desired label

Control Hazard: Branching

- ▶ Initial Solution: Stall until decision is made
 - Insert “**nop**” instructions (those that accomplish nothing, just take time) or hold up the fetch of the next instruction (for 2 cycles).
 - Drawback: branches take 3 clock cycles each (assuming comparator is put in ALU stage)
 - How many if statements (or loops) in your programs???

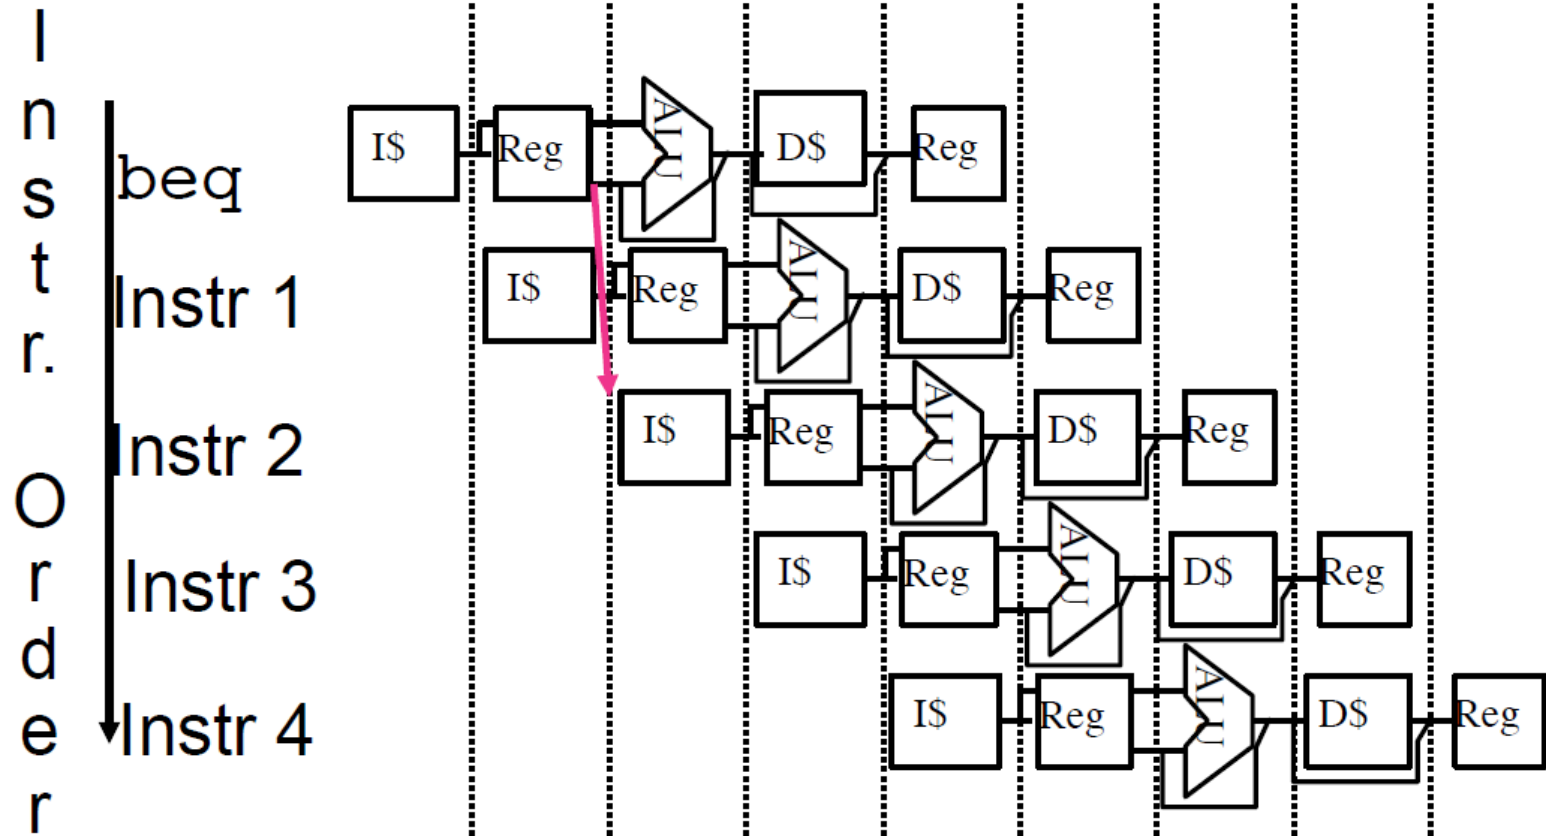
Control Hazard: Branching

► Optimization #1:

- Insert special branch comparator in Stage 2 (ID)
- As soon as instruction is decoded (Opcode identifies it as a branch), immediately make a decision and set the new value of the PC
- Benefit: since branch is complete in Stage 2, only one unnecessary instruction is fetched, so only one nop is needed
- Side Note: This means that branches are idle in Stages 3, 4 and 5.

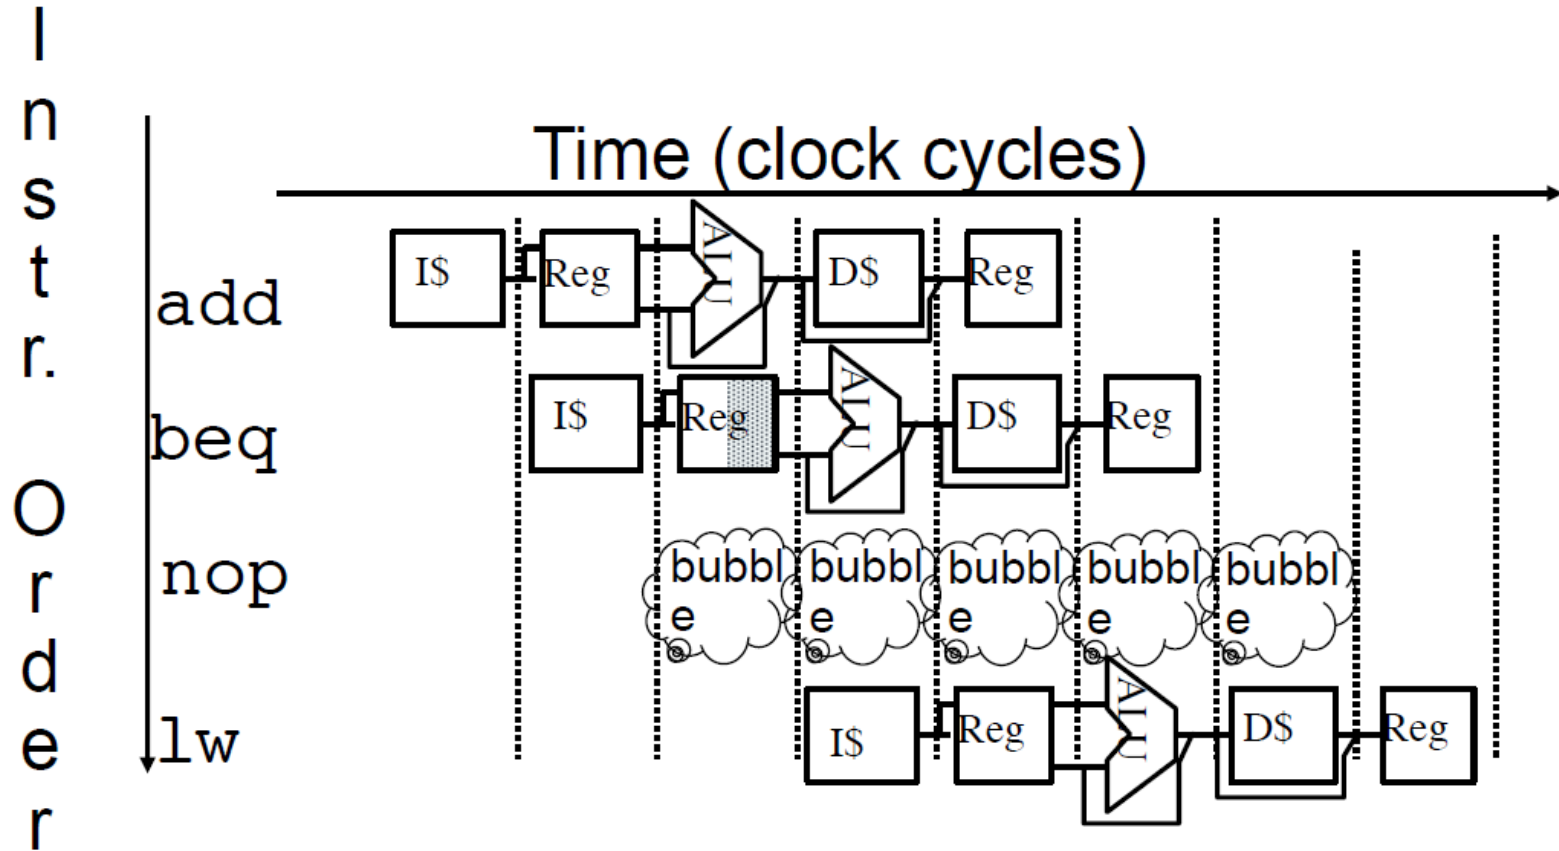
Control Hazard: Branching

Time (clock cycles)



- ▶ Branch comparator moved to Decode stage

Control Hazard: Branching



- ▶ User inserting nop instruction
- ▶ Impact: 2 clock cycles per branch instruction → waste

Control Hazard: Branching

- ▶ Optimization #2: Redefine branches
 - Old definition: if we take the branch, none of the instructions after the branch get executed by accident
 - **New definition**: whether or not we take the branch, the single instruction immediately following the branch gets executed (called the branch-delay slot)
- ▶ The term “**Delayed Branch**” means **we always execute instruction after branch**
- ▶ **This optimization is used with MIPS!**

Control Hazard: Branching

► Notes on Branch-Delay Slot

- Worst-Case Scenario: can always put a nop in the branch-delay slot
- Better Case: can find an instruction preceding the branch which can be placed in the branch-delay slot without affecting flow of the program
 - Video game after laundry → start video game while waiting for washer
 - Re-ordering instructions is a common method of speeding up programs (we can do this for Data Hazard from Load as well)
 - Compiler must be very smart in order to find instructions to do this
 - Usually can find such an instruction at least 50% of the time
 - Jumps also have a delay slot...

Nondelayed vs. Delayed Branch

Nondelayed Branch

or \$8, \$9, \$10

sub \$4, \$5, \$6

add \$1, \$2, \$3

beq \$1, \$4, Exit

xor \$10, \$1, \$11

Exit:

Delayed Branch

sub \$4, \$5, \$6

add \$1, \$2, \$3

beq \$1, \$4, Exit

or \$8, \$9, \$10

xor \$10, \$1, \$11

Exit:

Quiz

Assume 1 instr/clock, delayed branch, 5 stage pipeline, forwarding, interlock on unresolved load hazards (after 10^3 loops, so pipeline full)

How many pipeline stages (clock cycles) per loop iteration to execute this code?

```
Loop:      lw $t0, 0($s1)
           addu $t0, $t0, $s2
           sw $t0, 0($s1)
           addiu $s1, $s1, -4
           bne $s1, $zero, Loop
           addiu $s0, $t1, 4
```

1
2
3
4
5
6
7
8
9
10

Quiz

Assume 1 instr/clock, delayed branch, 5 stage pipeline, forwarding, interlock on unresolved load hazards (after 10^3 loops, so pipeline full)

How many pipeline stages (clock cycles) per loop iteration to execute this code?

2 (data hazard → stall)

Loop:

```
1 lw $t0, 0($s1)
3 addu $t0, $t0, $s2
4 sw $t0, 0($s1)
5 addiu $s1, $s1, -4
6 bne $s1, $zero, Loop
7 addiu $s0, $t1, 4
```

(delayed branch → nop)

1
2
3
4
5
6
7
8
9
10