

CSE 140

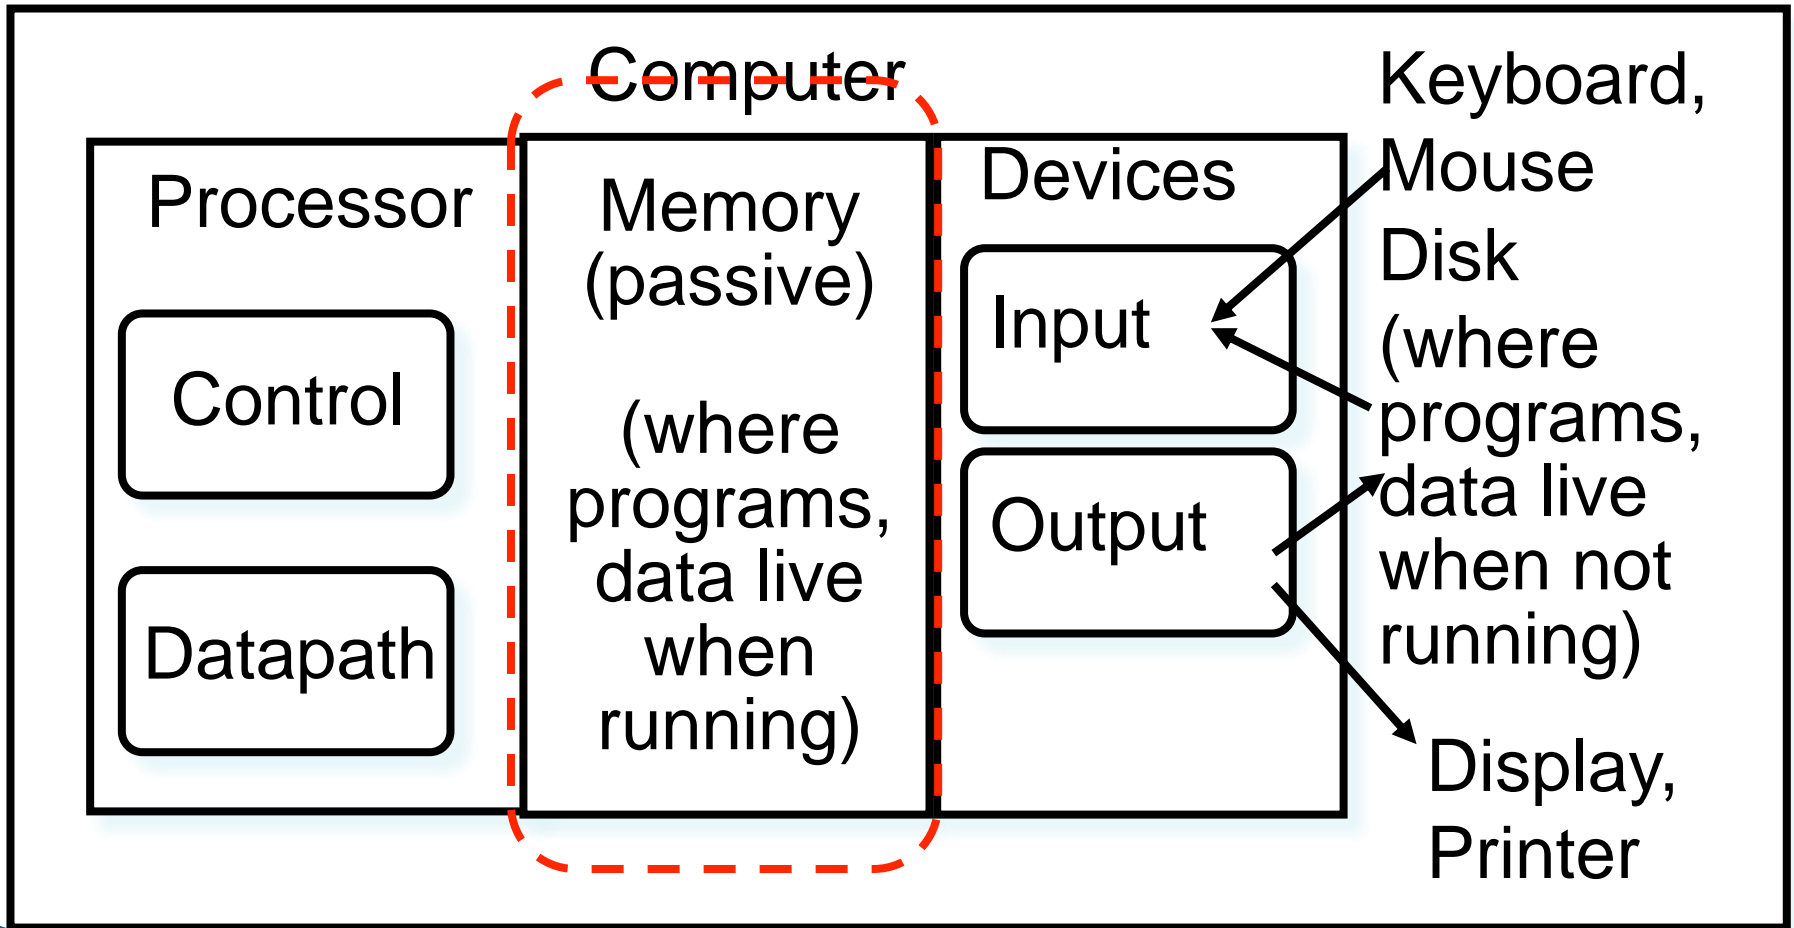
Computer Architecture

Lecture 6 – Memory (1)

Announcement

- ▶ Lab #2 this week
 - Due in one week
- ▶ Project #1
 - Due 10/11 (Monday) at 11:59pm
 - Start as soon as possible
- ▶ Reading assignment #3
 - Chapter 5.1 – 5.4
 - Do all **Participation Activities** in each section
 - Access through CatCourses
 - Due Thursday (9/19) at 11:59pm

Five Components of a Computer



Storage in Computer System

▶ Processor

- Holds data in register file (~100 Bytes)
- Registers accessed on nanosecond timescale

▶ Memory (we'll call it "main memory")

- More capacity than registers (~Gbytes)
- Access time ~50-100 ns
- Hundreds of clock cycles per memory access?!

▶ Disk

- HUGE capacity (virtually limitless)
- VERY slow: runs ~milliseconds
 - Samsung 960 Pro NVME M.2 (\$1,299, 2T)
 - 3,500MB/s -> 1.143 ms

Early Read-Only Memory Technologies





Punched cards, From early 1700s through Jacquard Loom, Babbage, and then IBM




Punched paper tape, instruction stream in Harvard Mk 1

Diode Matrix, EDSAC-2
 μ code store



IBM Card Capacitor ROS



IBM Balanced Capacitor
ROS


ROS – Read-only Storage

Early Read/Write Main Memory Technologies

Babbage, 1800s: Digits stored on mechanical wheels



Also, regenerative capacitor memory on Atanasoff-Berry computer, and rotating magnetic drum memory on IBM 650




Williams Tube, Manchester Mark 1, 1947

Mercury Delay Line, Univac 1, 1951



~1,000 words

MIT Whirlwind Core Memory



Core Memory

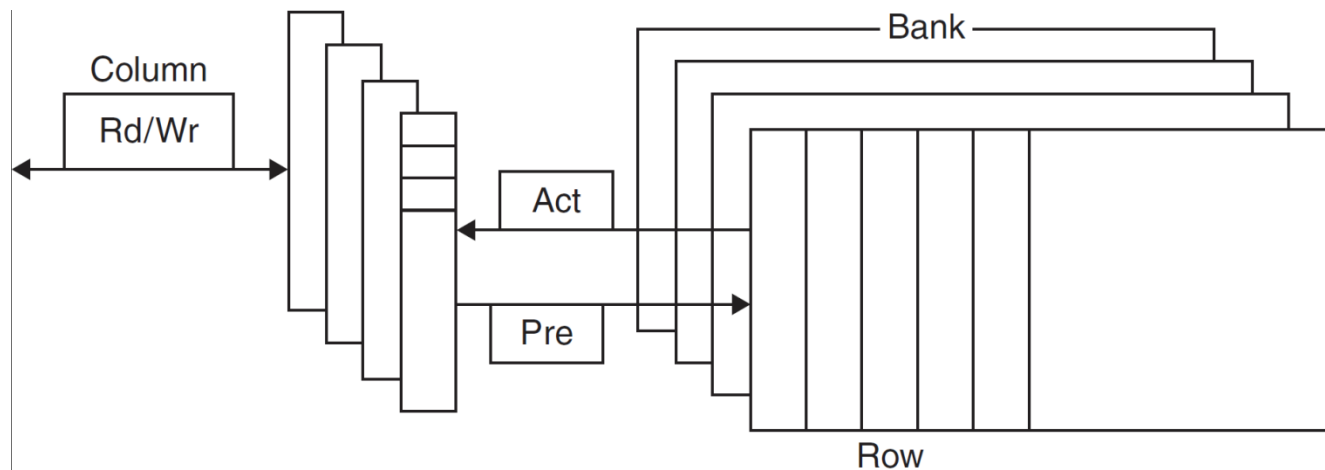
- ▶ Core memory was first large scale reliable main memory
 - invented by Forrester in late 40s/early 50s at MIT for Whirlwind project
- ▶ Bits stored as magnetization polarity on small ferrite cores threaded onto two-dimensional grid of wires
- ▶ Coincident current pulses on X and Y wires would write cell and also sense original state (destructive reads)
- ▶ Robust, non-volatile storage
- ▶ Used on space shuttle computers
- ▶ Cores threaded onto wires by hand (25 billion a year at peak production)
- ▶ Core access time $\sim 1\mu\text{s}$

Semiconductor Memory



- ▶ Semiconductor memory began to be competitive in early 1970s
 - Intel formed to exploit market for semiconductor memory
 - Early semiconductor memory was Static RAM (SRAM). SRAM cell internals similar to a latch (flip-flop).
- ▶ First commercial Dynamic RAM (DRAM) was Intel 1103
 - 1Kbit of storage on single chip
 - Charge on a capacitor used to hold value
- ▶ Semiconductor memory quickly replaced core in '70s

DRAM Technology


- ▶ Data stored as a charge in a capacitor
 - Single transistor used to access the charge
 - Must periodically be refreshed
 - Read contents and write back
 - Performed on a DRAM “row”



One-Transistor Dynamic RAM [Dennard, IBM]

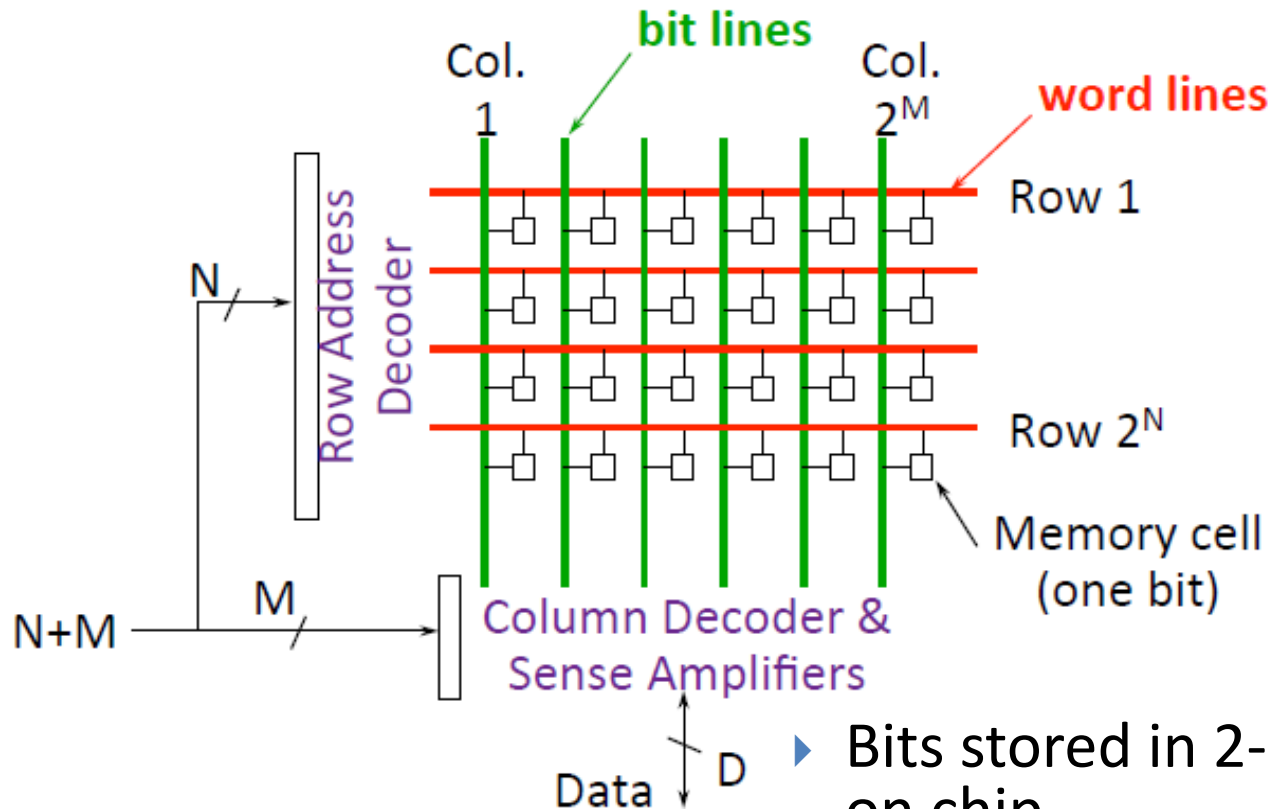


Modern DRAM Structure



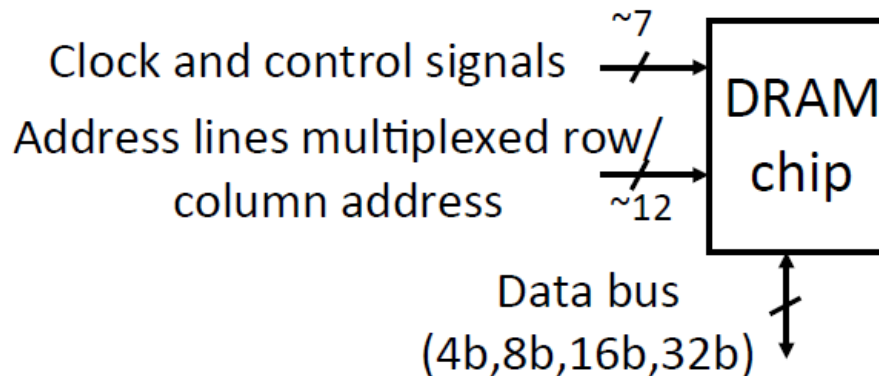
[Samsung, sub-70nm DRAM, 2004]

DRAM Architecture



- ▶ Bits stored in 2-dimensional arrays on chip
- ▶ Modern chips have around 4-8 logical banks on each chip
 - Each logical bank physically implemented as many smaller arrays


DRAM Packaging (Laptops/Desktops/Servers)



- ▶ DIMM (Dual Inline Memory Module) contains multiple chips with clock/control/address signals connected in parallel (sometimes need buffers to drive signals to all chips)
- ▶ Data pins work together to return wide word (e.g., 64-bit data bus using 16x4-bit parts)



72-pin SO DIMM



168-pin DIMM

DRAM Packaging, Mobile Devices



[Apple A4 package on circuit board]



Two stacked
DRAM die
Processor plus
logic die


[Apple A4 package cross-section, iFixit 2010]

Advanced DRAM Organization

- ▶ Bits in a DRAM are organized as a rectangular array
 - DRAM accesses an entire row
 - Burst mode: supply successive words from a row with reduced latency
- ▶ Double data rate (DDR) DRAM
 - Transfer on rising and falling clock edges
- ▶ Quad data rate (QDR) DRAM
 - Separate DDR inputs and outputs

DRAM Generations

Year	Capacity	\$/GB
1980	64Kbit	\$1500000
1983	256Kbit	\$500000
1985	1Mbit	\$200000
1989	4Mbit	\$50000
1992	16Mbit	\$15000
1996	64Mbit	\$10000
1998	128Mbit	\$4000
2000	256Mbit	\$1000
2004	512Mbit	\$250
2007	1Gbit	\$50

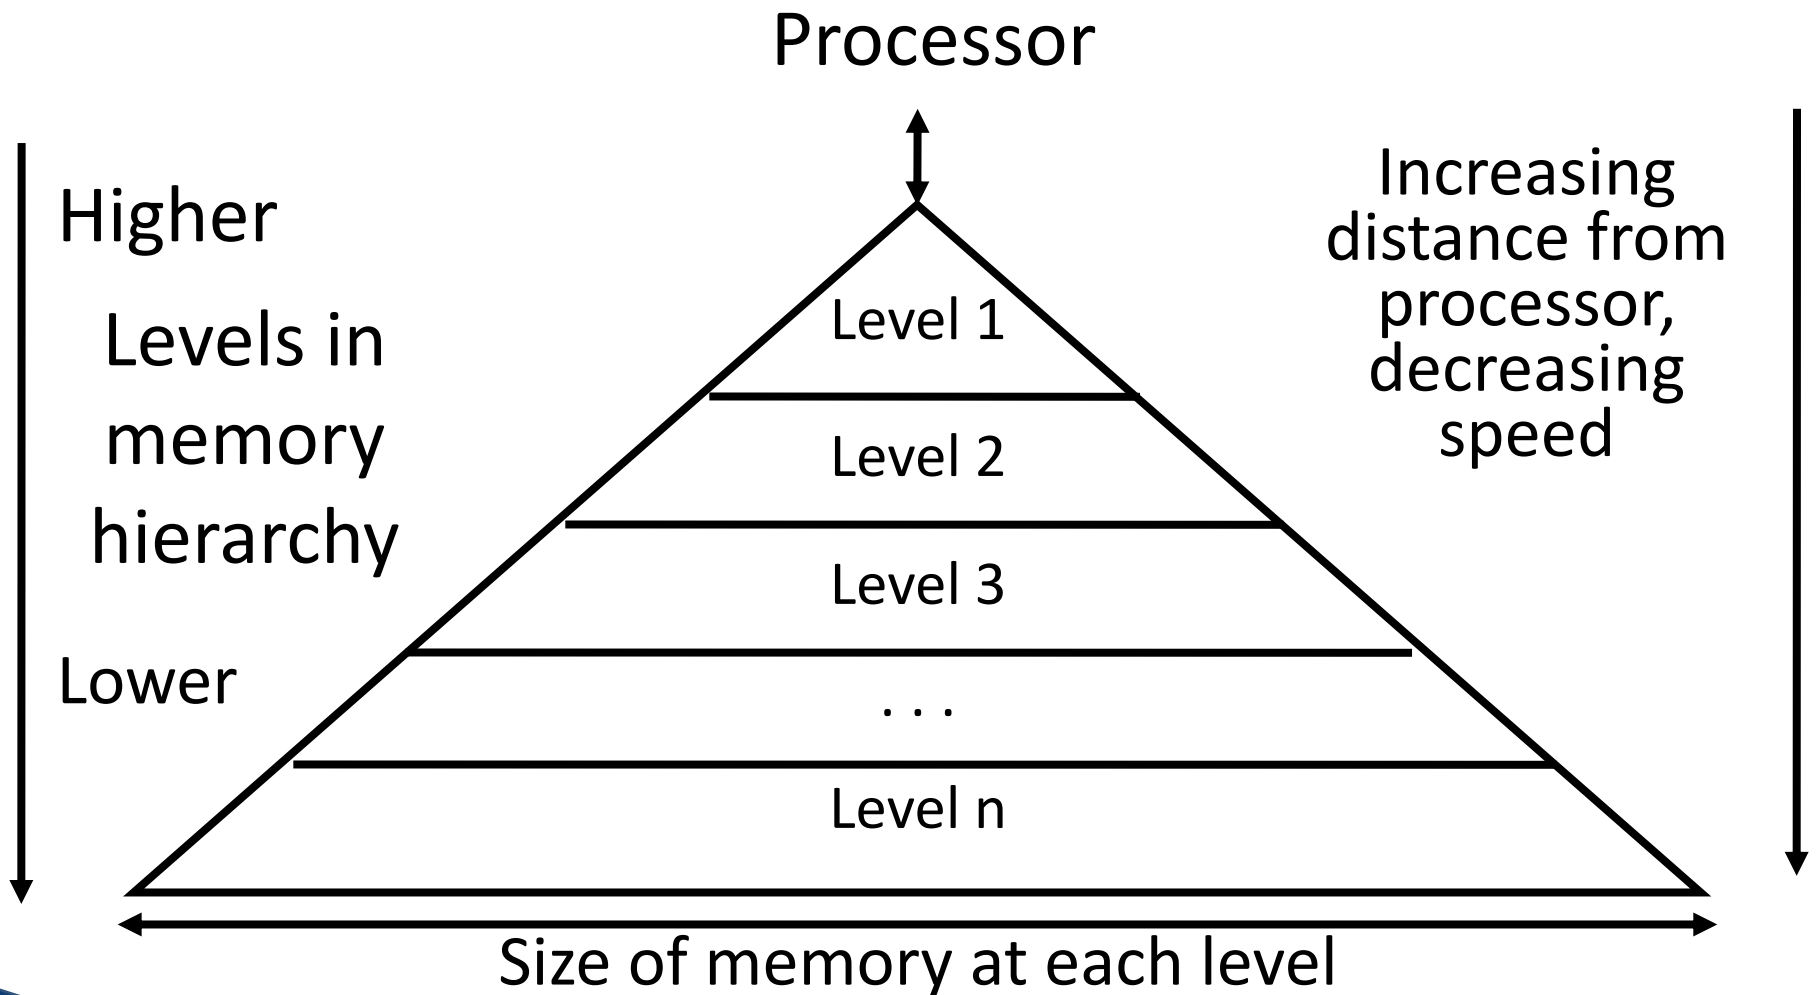


Trac: Random access time
Tcac: Column access time

Memory Caching

- ▶ Processor and memory speed mismatch leads us to add a new level: a memory **cache**
- ▶ Implemented with same integrated circuit processing technology as processor, integrated on-chip: faster but more expensive than DRAM memory
- ▶ **Cache is a copy of a subset of main memory**
- ▶ Modern processors have separate caches for instructions and data, as well as several levels of caches implemented in different sizes (why?)
 - Often use \$ (“cash”) to abbreviate cache, e.g. D\$ = Data Cache, I\$ = Instruction Cache

Memory Hierarchy



As we move to deeper levels the latency goes up and price per bit goes down. Why?

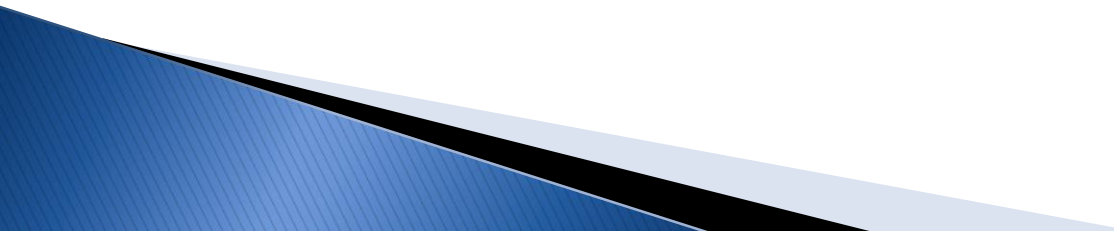
Memory Hierarchy

- ▶ If level closer to Processor, it is:
 - Smaller
 - Faster
 - More expensive
 - Subset of lower levels (contains most recently used data)
- ▶ Lowest Level (usually disk) contains all available data (does it go beyond the disk?)
- ▶ Memory Hierarchy presents the processor with the illusion of a very large & fast memory

Memory Hierarchy Basis

- ▶ Cache contains copies of data in memory that are being used.
- ▶ Memory contains copies of data on disk that are being used.
- ▶ Caches work on the principles of **temporal** and **spatial** locality.
 - **Temporal Locality**: if we use it now, chances are we'll want to use it again soon.
 - **Spatial Locality**: if we use a piece of memory, chances are we'll use the neighboring pieces soon.

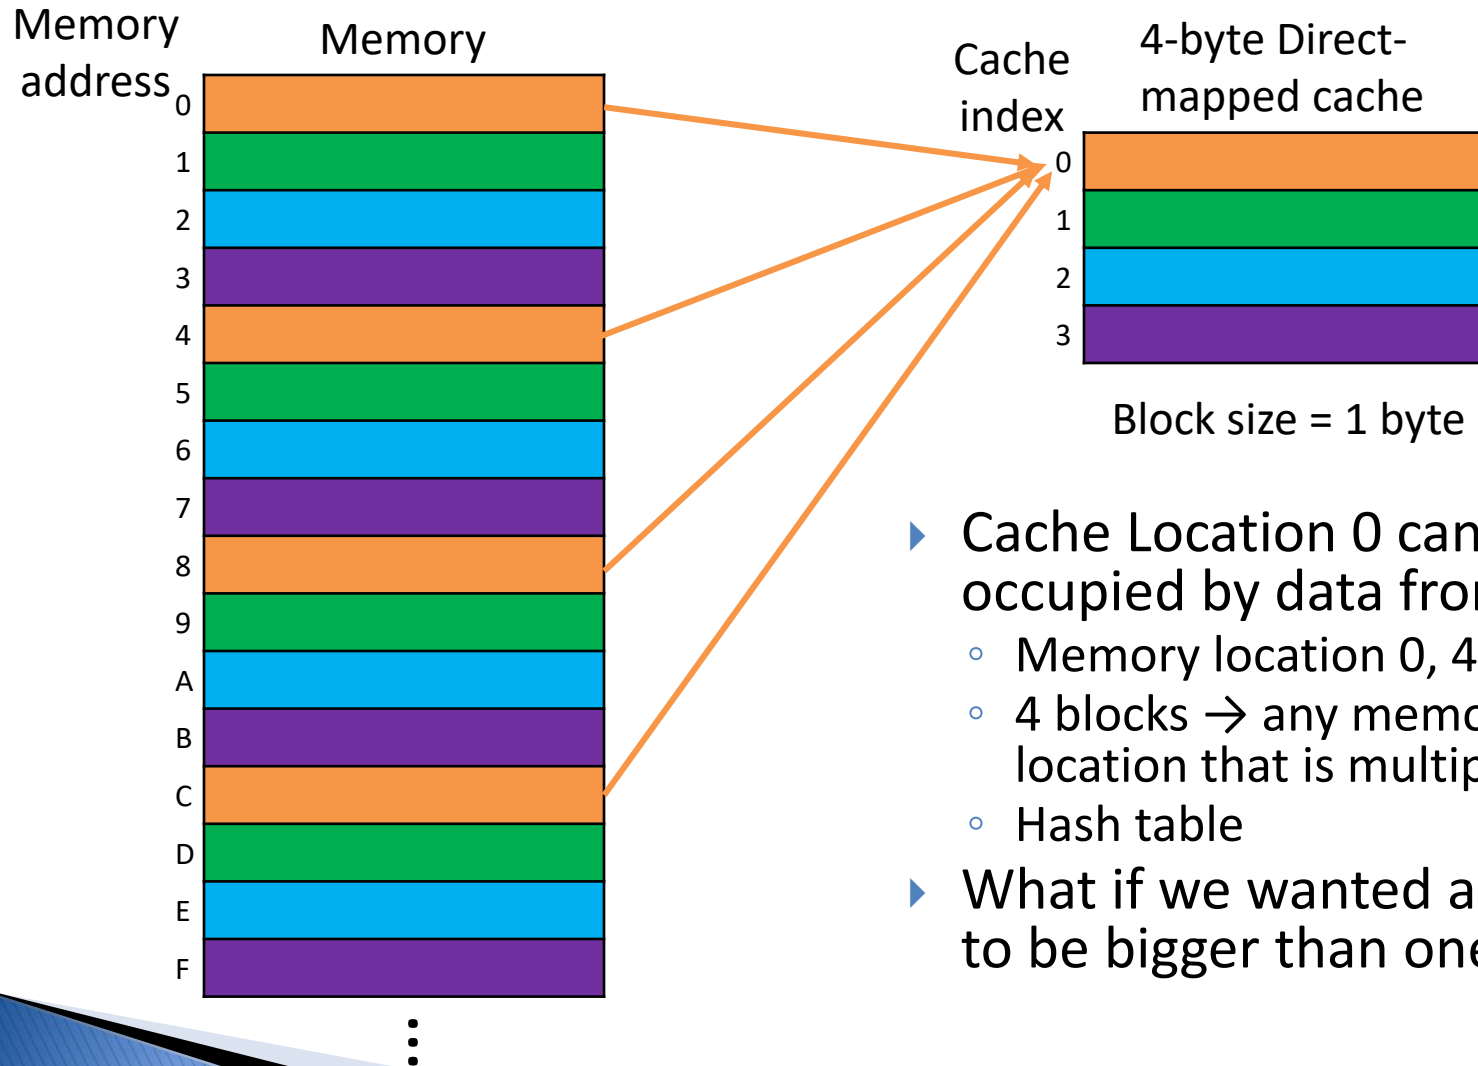
Cache Design Questions

- ▶ How do we organize cache?
 - ▶ Where does each memory address map to?
 - Remember that cache is a subset of memory, so multiple memory addresses can map to the same cache location.
 - ▶ How do we know which elements are in cache?
 - ▶ How do we quickly locate them?
- 

Direct-Mapped Cache (1/4)

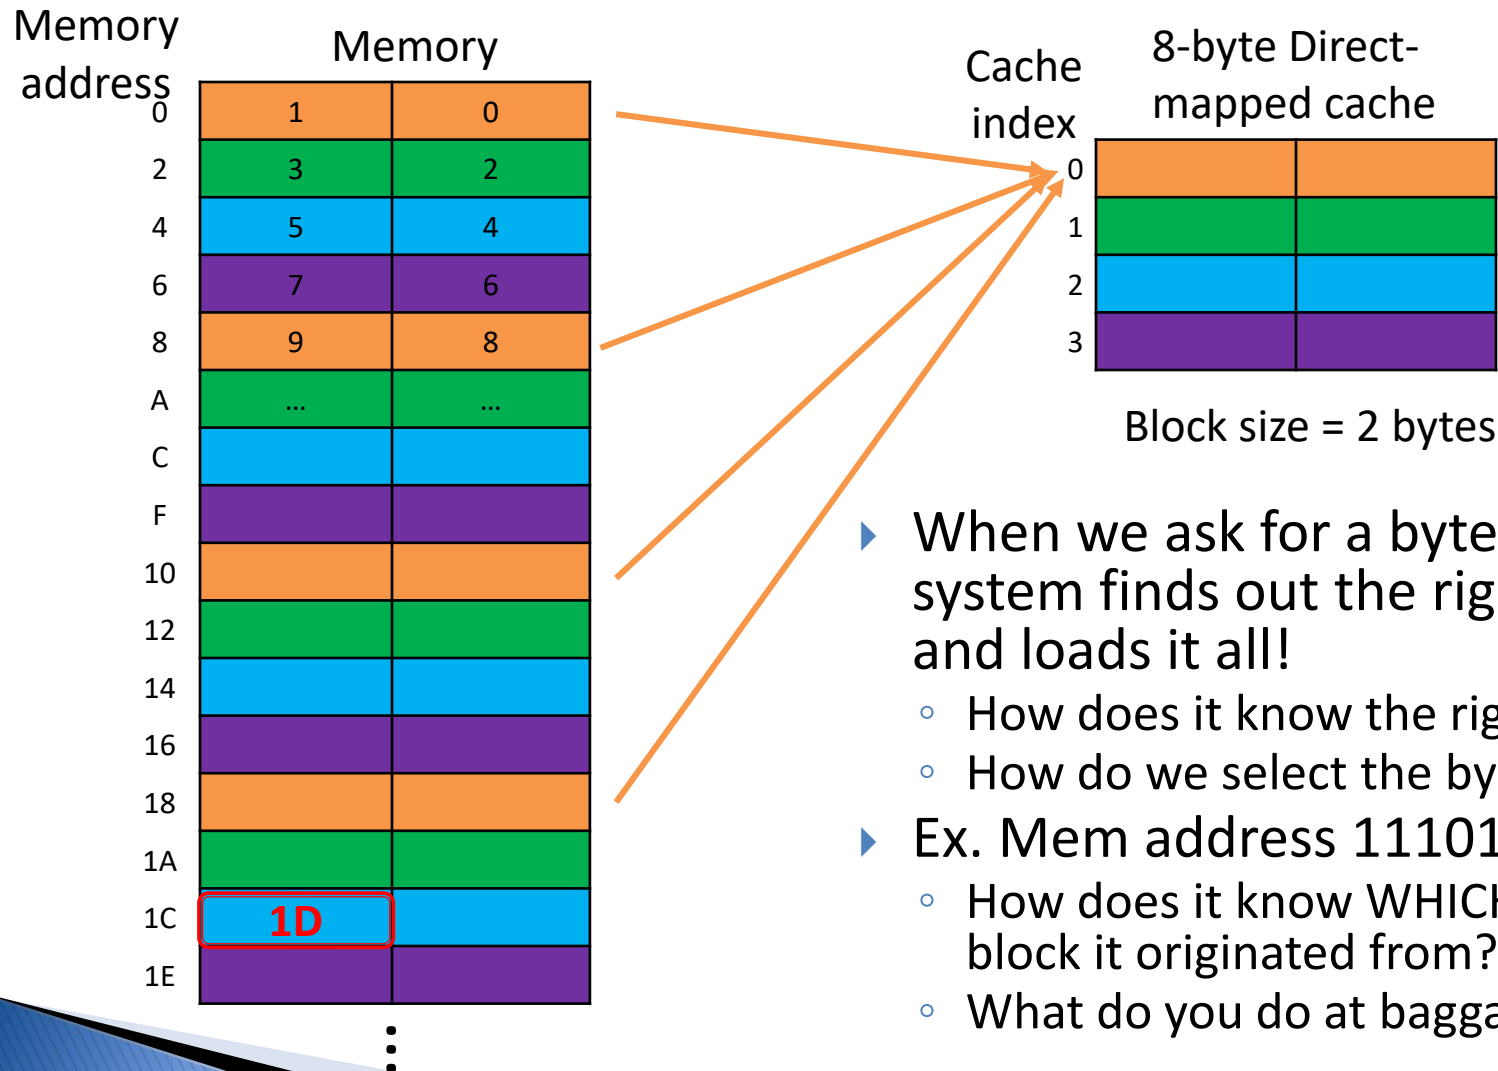
- ▶ In a direct-mapped cache, each memory address is associated with one possible block within the cache
 - Therefore, we only need to look in a single location in the cache for the data if it exists in the cache
 - Block is the unit of transfer between cache and memory

Direct-Mapped Cache (2/4)

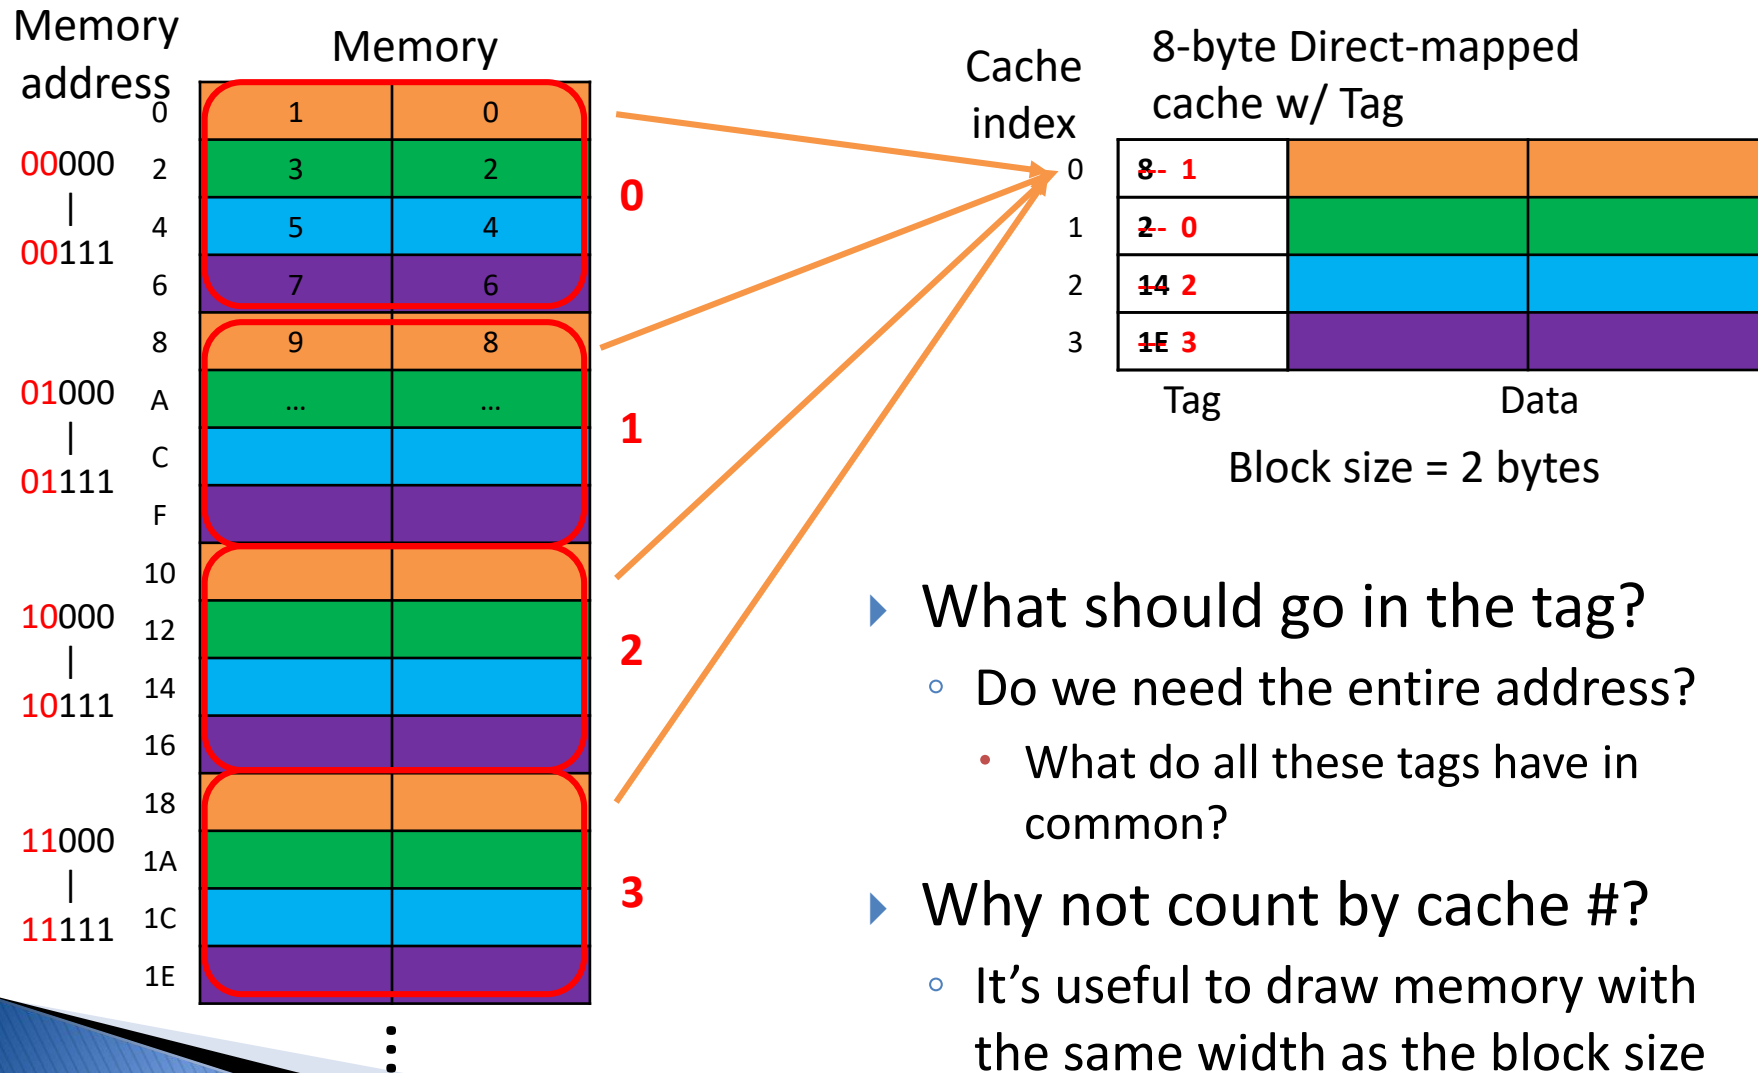


- ▶ Cache Location 0 can be occupied by data from:
 - Memory location 0, 4, 8, ...
 - 4 blocks → any memory location that is multiple of 4
 - Hash table
- ▶ What if we wanted a block to be bigger than one byte?

Direct-Mapped Cache (3/4)



Direct-Mapped Cache (4/4)



Issues with Direct-Mapped

- ▶ Since multiple memory addresses map to same cache index, how do we tell which one is in there?
- ▶ What if we have a block size > 1 byte?
 - Answer: divide memory address into three fields

t t t t t t t t t t t t t t t t	i i i i i i i i i i	o o o o
---------------------------------	---------------------	---------

Tag to check if it has the correct block

Index to select block

Byte
offset
within
block

Direct-Mapped Cache Terminology

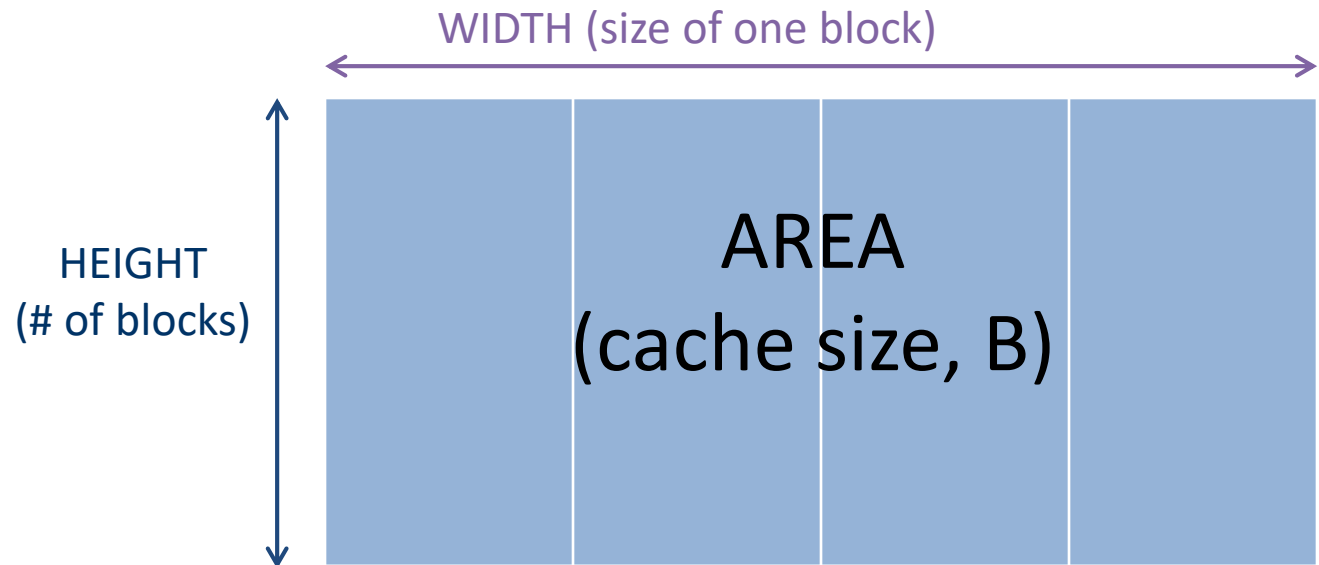
- ▶ All fields are read as unsigned integers
- ▶ **Index**
 - specifies the cache index (which “row” or block of the cache we should look in)
- ▶ **Offset**
 - once we’ve found correct block, specifies which byte within the block we want
- ▶ **Tag**
 - the remaining bits after offset and index are determined; these are used to distinguish between all the memory addresses that map to the same location



AREA (cache size, B)

= HEIGHT (# of blocks) * WIDTH (size of one block)

Tag	Index	Offset
-----	-------	--------



Direct-Mapped Cache Example (1/3)

- ▶ Suppose we have a 8B of data in a direct-mapped cache with 2 byte blocks
 - Sound familiar?
- ▶ Determine the number of bits in the tag, index and offset fields if we're using a 32-bit architecture
- ▶ **Offset**
 - need to specify correct byte within a block
 - block contains 2 bytes
 - $= 2^1$ bytes
 - need **1 bit** to specify correct byte

Direct-Mapped Cache Example (2/3)

- ▶ **Index:** (~index to an “array of blocks”)
 - need to specify correct number of block in cache
 - Cache contains 8 B = 2^3 bytes
 - block contains 2 B = 2^1 bytes
 - # blocks/cache
 - = $\frac{\text{bytes/cache}}{\text{bytes/block}}$
 - = $\frac{2^3 \text{ bytes/cache}}{2^1 \text{ bytes/block}}$
 - = 2^2 blocks/cache
 - need **2 bits** to specify this many blocks

Direct-Mapped Cache Example (3/3)

- ▶ **Tag:** use remaining bits as tag
 - tag length = addr length – offset – index
= 32 - 1 - 2 bits
= 29 bits
 - so tag is leftmost **29 bits** of memory address
- ▶ Why not full 32 bit address as tag?
 - Index must be same for every address within a block, so it's redundant in tag check, thus can leave off to save memory
- ▶ Each row has Valid and Dirty bit

Caching Terminology

- ▶ When reading memory, 3 things can happen:
 - cache hit:
 - cache block is valid and contains proper address, so read desired word
 - cache miss:
 - nothing in cache at appropriate block, so fetch from memory
 - cache miss, block replacement:
 - wrong data is in cache at appropriate block, so discard it and fetch desired data from memory (cache always copy)

Accessing data in a direct-mapped cache

- ▶ Ex.: 16KB of data, direct-mapped, 4 word blocks (16B)
 - Can you work out height, width, area?
- ▶ Read 4 addresses (Hex)

1. 0x00000014
2. 0x0000001C
3. 0x00000034
4. 0x00008014

...	...
00000010	a
00000014	b
00000018	c
0000001C	d
...	...
00000030	e
00000034	f
00000038	g
0000003C	h
...	...
00008010	i
00008014	j
00008018	k
0000801C	l
...	...

Memory

Accessing data in a direct-mapped cache

- ▶ 4 Addresses:
 - 0x00000014, 0x0000001C, 0x00000034, 0x00008014
- ▶ 4 Addresses divided (for convenience) into **Tag**, **Index**, **Byte Offset** fields

Tag	Index	Offset
000000000000000000000000	00000000001	0100
000000000000000000000000	00000000001	1100
000000000000000000000000	00000000011	0100
000000000000000000000010	00000000001	0100

16 KB Direct Mapped Cache, 16B blocks

- ▶ **Valid bit:** determines whether anything is stored in that row (start with all entries invalid)

Index	Valid	Tag	0xC-F	0x8-B	0x4-7	0x0-3
0	0					
1	0					
2	0					
3	0					
4	0					
5	0					
6	0					
7	0					
...				...		
1022	0					
1023	0					

Read 0x00000014

Tag

Index

Offset

000000000000000000000000 000000000001 0100

Index	Valid	Tag	0xC-F	0x8-B	0x4-7	0x0-3
0	0					
1	0					
2	0					
3	0					
4	0					
5	0					
6	0					
7	0					
...				...		
1022	0					
1023	0					

Read block 1

Tag

Index

Offset

000000000000000000000000 0000000001 0100

Index	Valid	Tag	0xC-F	0x8-B	0x4-7	0x0-3
0	0					
1	0					
2	0					
3	0					
4	0					
5	0					
6	0					
7	0					
...				...		
1022	0					
1023	0					

No valid data

Tag

Index

Offset

000000000000000000000000 0000000001 0100

Index	Valid	Tag	0xC-F	0x8-B	0x4-7	0x0-3
0	0					
1	0					
2	0					
3	0					
4	0					
5	0					
6	0					
7	0					
...				...		
1022	0					
1023	0					

Load data into cache, set tag and valid

Tag

Index

Offset

00000000000000000000

00000000001

0100

Index	Valid	Tag	0xC-F	0x8-B	0x4-7	0x0-3
0	0					
1	1	0	d	c	b	a
2	0					
3	0					
4	0					
5	0					
6	0					
7	0					
...				...		
1022	0					
1023	0					

Read from cache at offset

Tag

Index

Offset

000000000000000000000000 000000000001 0100

Index	Valid	Tag	0xC-F	0x8-B	0x4-7	0x0-3
0	0					
1	1	0	d	c	b	a
2	0					
3	0					
4	0					
5	0					
6	0					
7	0					
...				...		
1022	0					
1023	0					

Return word b

Read 0x0000001C

Tag

Index

Offset

000000000000000000000000 00000000001 1100

Index	Valid	Tag	0xC-F	0x8-B	0x4-7	0x0-3
0	0					
1	1	0	d	c	b	a
2	0					
3	0					
4	0					
5	0					
6	0					
7	0					
...				...		
1022	0					
1023	0					

Index is valid

Tag

Index

Offset

000000000000000000000000 0000000001 1100

Index	Valid	Tag	0xC-F	0x8-B	0x4-7	0x0-3
0	0					
1	1	0	d	c	b	a
2	0					
3	0					
4	0					
5	0					
6	0					
7	0					
...				...		
1022	0					
1023	0					

Tag matches

							Tag	Index	Offset
							00000000000000000000	000000000001	1100
Index	Valid	Tag	0xC-F	0x8-B	0x4-7	0x0-3			
0	0								
1	1	0	d	c	b	a			
2	0								
3	0								
4	0								
5	0								
6	0								
7	0								
...									
1022	0								
1023	0								

Read from cache at offset

Tag

Index

Offset

000000000000000000000000 000000000001 1100

Index	Valid	Tag	0xC-F	0x8-B	0x4-7	0x0-3
0	0					
1	1	0	d	c	b	a
2	0					
3	0					
4	0					
5	0					
6	0					
7	0					
...				...		
1022	0					
1023	0					

Return word d

Read 0x00000034

Tag

Index

Offset

000000000000000000000000 00000000011 0100

Index	Valid	Tag	0xC-F	0x8-B	0x4-7	0x0-3
0	0					
1	1	0	d	c	b	a
2	0					
3	0					
4	0					
5	0					
6	0					
7	0					
...				...		
1022	0					
1023	0					

Read block 3

Tag

Index

Offset

000000000000000000000000 0000000011 0100

Index	Valid	Tag	0xC-F	0x8-B	0x4-7	0x0-3
0	0					
1	1	0	d	c	b	a
2	0					
3	0					
4	0					
5	0					
6	0					
7	0					
...				...		
1022	0					
1023	0					

No valid data

Tag

Index

Offset

000000000000000000000000 0000000011 0100

Index	Valid	Tag	0xC-F	0x8-B	0x4-7	0x0-3
0	0					
1	1	0	d	c	b	a
2	0					
3	0					
4	0					
5	0					
6	0					
7	0					
...				...		
1022	0					
1023	0					

Load data, read cache at offset

Tag				Index	Offset	
000000000000000000000000				000000000011	0100	
Index	Valid	Tag	0xC-F	0x8-B	0x4-7	0x0-3
0	0					
1	1	0	d	c	b	a
2	0					
3	1	0	h	g	f	e
4	0					
5	0					
6	0					
7	0					
...				...		
1022	0					
1023	0					

Return word f

Read 0x00008014

Tag

Index

Offset

000000000000000000000010 000000000001 0100

Index	Valid	Tag	0xC-F	0x8-B	0x4-7	0x0-3
0	0					
1	1	0	d	c	b	a
2	0					
3	1	0	h	g	f	e
4	0					
5	0					
6	0					
7	0					
...				...		
1022	0					
1023	0					

Read block 1, valid data

Tag

Index

Offset

0000000000000000000010 0000000001 0100

Index	Valid	Tag	0xC-F	0x8-B	0x4-7	0x0-3
0	0					
1	1	0	d	c	b	a
2	0					
3	1	0	h	g	f	e
4	0					
5	0					
6	0					
7	0					
...				...		
1022	0					
1023	0					

Tag not matched!

Tag

Index

Offset

000000000000000000010

000000000001

0100

Index	Valid	Tag	0xC-F	0x8-B	0x4-7	0x0-3
0	0					
1	1	0	d	c	b	a
2	0					
3	1	0	h	g	f	e
4	0					
5	0					
6	0					
7	0					
...				...		
1022	0					
1023	0					

Miss: replace block 1

Tag

Index

Offset

00000000000000000010

00000000001

0100

Index	Valid	Tag	0xC-F	0x8-B	0x4-7	0x0-3
0	0					
1	1	2	l	k	j	i
2	0					
3	1	0	h	g	f	e
4	0					
5	0					
6	0					
7	0					
...				...		
1022	0					
1023	0					

Read cache at offset

Tag

Index

Offset

0000000000000000000010 000000000001 0100

Index	Valid	Tag	0xC-F	0x8-B	0x4-7	0x0-3
0	0					
1	1	2	l	k	j	i
2	0					
3	1	0	h	g	f	e
4	0					
5	0					
6	0					
7	0					
...				...		
1022	0					
1023	0					

Return word j

Do these examples yourself

- ▶ Read address 0x00000030
- ▶ Read address 0x0000001C
- ▶ Cache: Hit, Miss, or Miss with replace?
- ▶ Returned values: a, b, c, ..., k, l?

Index	Valid	Tag	0xC-F	0x8-B	0x4-7	0x0-3
0	0					
1	1	2	l	k	j	i
2	0					
3	1	0	h	g	f	e
4	0					
5	0					
6	0					
7	0					

Answers

- ▶ 0x00000030: hit
 - Index = 3, Tag matches
 - Offset = 0, returned value = e
- ▶ 0x0000001c: miss
 - Index = 1, Tag mismatch, so replace from memory
 - Offset = 0xc, value = d

Summary

- ▶ Mechanism for transparent movement of data among levels of a storage hierarchy
 - set of address/value bindings
 - address \Rightarrow index to set of candidates
 - compare desired address with tag
 - service hit or miss
 - load new block and binding on miss

